

Hive-ODCI - Users Guide

Hive-ODCI is an [Oracle Data Cartridge Interface](#) for dynamically accessing Hadoop/Hive data-stores through an Oracle 12c database. In other words Hive-ODCI makes Hadoop/Hive tables accessible as first-class, native, objects directly using PL/SQL, SQL, VIEWS, DML, DDL, etc.... in an Oracle 12c database.

Author

Metasystem Technologies Inc. (MTI)
www.mtihq.com

Nicholas Van Wyen
nvanwyen@mtihq.com

License

**Hive-ODCI - Copyright (C) 2006-2016 Metasystems Technologies Inc. (MTI)
Nicholas Van Wyen**

This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2.1 of the License, or (at your option) any later version.

This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the

Free Software Foundation, Inc.
59 Temple Place, Suite 330,
Boston, MA 02111-1307 USA

Releases

All releases can be found on Github
<https://github.com/nvanwyen/hive-odci/releases>, along with the [latest release](#).

The project home is publicly available on Github at
<https://github.com/nvanwyen/hive-odci>

Installation and Removal

See `INSTALL.md` for instructions

Concepts

Hive-ODCI is a pass-through interface allowing SQL access from within an Oracle RDBMS to information retained in an external Hive/Hadoop data-store. Hive-ODCI provides PL/SQL interfaces using ODCI to accomplish this functionality, making the access viable as native object in Oracle.

Hive-ODCI is accessed via the `HIVE` schema, and controlled through RBAC permissions in Oracle. A client (user or application) is granted privileges through the `HIVE_USER` role or direct system privileges by the DBA.

The client uses the PL/SQL objects to query, or execute DML/DDDL in the remote Hive datastore, using the `HIVE_T` object type or one of the PL/SQL packages.

The `HIVE` schema installs Java classes in the database, which perform the JDBC execution on the clients behalf, and streams `PIPLINED` results back through the same interface.

The client can dictate most levels of functionality at run-time, with predefined session data, bind-variables, etc... as customization for each call.

Take for example the following concepts. The Hive/Hadoop data-store is remotely accessible on a separate server. The Hive-ODCI Java classes access the remote data via the JDBC Driver loaded during installation.

The client accesses the Hive-ODCI interface using the PL/SQL objects provided and/or a first-class `VIEW`, controlled by RBAC, for a 2-way avenue for data.

Our tables looks like this ...

```
SQL> desc SCOTT.USER_LOG
```

Name	Null?	Type
STAMP	NOT NULL	DATE
ACCOUNT		VARCHAR2(30)
MESSAGE		VARCHAR2(4000)

Because we have tons of room available in our Hadoop cluster we decide that we want to move the data there, so it can be indexed and searched. But hold on, we have a problem, we still have an application that reads the table and creates reports for upper management and they are not going to change their application to read from 2 different places using 2 different methods (oh, what to do).

Hive-ODCI to the rescue

On top of inserting, updating and deleting capabilities to the table, the application also contains PL/SQL to create the reports and has a VIEW used in displaying the monthly metrics.

They look something like this ...

```
--
procedure user_log_report( p_report out xmltype ) is
begin

    for rec in ( select account,
                      message
                  from scott.user_log
                  order by account ) loop

        if ( rec.stamp > sysdate - 90 ) then

            p_report := ... -- do something special

        else

            if ( rec.account = user ) then

                p_report := ... -- if current user then ...

            end if;

        end if;

        if ( rec.message like '%ABC%' ) then

            p_report := ... -- write a particular format

        else

            p_report := ... -- write another format

        end if;

    end loop;

end;
```

```

        end if;

    end loop;

end user_log_report;

--
view user_log_monthly
as
select stamp,
       account,
       message
  from scott.user_log
 where stamp between sysdate - 30
            and sysdate;

```

Assume we have already moved our data over to Hadoop and created a Hive table, of the same name. Our remote table looks like this ...

```

$ beeline -u jdbc:hive2://hive.corp.com:10000 \
-n oracle \
-w welcome1.passwd

0: jdbc:hive2://localhost:10000> desc user_log;
+-----+-----+-----+
| col_name | data_type | comment |
+-----+-----+-----+
| stamp    | date      |         |
| account  | string    |         |
| message  | string    |         |
+-----+-----+-----+
3 rows selected (0.17 seconds)

```

Hive-ODCI configuration

Now that we have our table in Hadoop/Hive and we can connect via beeline, and see it there, let's setup Hive-ODCI as an access point to that table.

Because all of our clients will be accessing the same Hadoop/Hive data-store, we can setup a common connection strategy for Hive-ODCI via the parameters

```

dbms_hive.param( 'hive_jdbc_url', 'jdbc:hive2://hive.corp.com:10000' );
dbms_hive.param( 'hive_jdbc_url.1', 'user=oracle' );
dbms_hive.param( 'hive_jdbc_url.2', 'password=welcome1' );

```

If we have clients which need a different connection strategy or use different parameters, they can change these at the session level to meet their needs specifically, but for now we'll assume everyone is using the same thing.

Hive-ODCI object creation

As denoted above we have PL/SQL code and a VIEW which is now invalid because the `SCOTT.USER_LOG` table no longer exists.

So, let's put it back using Hive-ODCI. First let's create a new VIEW which replaces the table.

```
--  
grant execute on hive_q to scott;  
  
--  
create or replace view scott.user_log  
(  
    stamp,  
    account,  
    message  
)  
as  
select *  
from table( hive_q( q'[ select stamp,  
                           account,  
                           message  
                           from user_log  
                           order by stamp ]' ) )  
  
/
```

Whew, that was easy. We now have a column-by-column replacement of the old table with a newly created remote table, that provides the same data types as before.

```
SQL> desc SCOTT.USER_LOG
```

Name	Null?	Type
STAMP	NOT NULL	DATE
ACCOUNT		VARCHAR2(4000)
MESSAGE		VARCHAR2(4000)

Let's see if it worked ...

```
SQL> alter procedure scott.user_log_report compile;  
  
Procedure altered.
```

Excellent, so let's take a look at replacing the view. In this case, we only care about data from the last 30 days, no need to make Hadoop/Hive do more work than it has to. Let's use bindings to restrict the data at the Hadoop/Hive layer instead of at the Oracle layer.

```
--  
grant execute on hive_bind to scott;  
grant execute on hive_binds to scott;
```

```
--
create or replace view scott.user_log_monthly
(
    stamp,
    account,
    message
)
as
select *
from table( hive_q( q'[ select stamp,
                        account,
                        message
                        from user_log
                        where stamp between ? and ? ]',
hive_binds( hive_bind( to_char( sysdate - 30,
                                'yyyy-mm-dd' ),
                        1 /* type_date */,
                        1 /* ref_in */ ),
hive_bind( to_char( sysdate, ,
                    'yyyy-mm-dd' ),
            1 /* type_date */,
            1 /* ref_in */ ) ) )
/
```

Let's break this one down, as it is a little more complex than the other. The `hive_q` function takes 3 parameters, 2 of which are defaulted to `NULL`. The first is a `hive_binds` object which is simply an array of `hive_bind` objects or individual bind data. The second parameter is the `hive_session` (e.g. URL, User, ...), but since we are using a common connection strategy we will ignore this argument, not pass it in, let it continue to be `NULL`.

The `hive_bind` object also takes 3 arguments, the first is the data to be used as the bind. The second is the type of data to be bound (e.g. string, date, number, etc...). And the third is the scope of the bind, most useful for DML operations (e.g. IN, OUT, IN/OUT).

In our case we only care about scope references of IN, hence the `1 /* ref_in */`. Since both bind operators are `DATE` variables they both have `1 /* type_date */`.

The actual bound data is based on `SYSDATE`, but since we don't know up front the `NLS_DATE_FORMAT` setting for the session, we simply guarantee the format by wrapping it in a `TO_CHAR()` function and specify the format Hadoop/Hive is expecting.

The Oracle RBAC controls dictate who can SELECT from the VIEWS, just as before. So, if you have custom roles which need access to the VIEWS, they can be granted access in the same way.

```
grant select on scott.user_log to my_app_role;
grant select on scott.user_log_monthly to public;
```

Now wala! We have a VIEW that will retrieve the last 30 days worth of data from Hadoop/Hive.

Almost there

We now have replacements for both the PL/SQL Reporting and the monthly VIEW , but we're just not yet at a 100%, so what have we forgot?

You guessed it; DML. Our application is still putting data in, modifying it and removing it so we need to support that too. No problem. We can use an `INSTEAD OF` trigger and send our `INSERT`, `UPDATE` and `DELETE` commands to Hadoop/Hive using Hive-ODCI.

```

create or replace trigger scott.user_log_dml

instead of delete or insert or update on scott.user_log
for each row

declare

cmd varchar2( 4000 );
bnd hive_binds := hive_binds();

begin

if ( inserting ) then

cmd := q'[ insert into user_log
              ( stamp, account, message )
            values
              ( ?, ?, ? ) ]';

bnd.extend;
bnd( bnd.count ) := hive_bind( to_char( :new.stamp,
                                         'yyyy-mm-dd' ),
                               hive_binding.type_date,
                               hive_binding.ref_in );

bnd.extend;
bnd( bnd.count ) := hive_bind( :new.account,
                               hive_binding.type_string,
                               hive_binding.ref_in );

bnd.extend;
bnd( bnd.count ) := hive_bind( :new.message,
                               hive_binding.type_string,
                               hive_binding.ref_in );

elsif ( updating ) then

cmd := q'[ update user_log
              set account = ?,
                message = ?
            where stamp = ? ]';

bnd.extend;
bnd( bnd.count ) := hive_bind( :new.account,
                               hive_binding.type_string,
                               hive_binding.ref_in );

```



```

        bnd.extend;
        bnd( bnd.count ) := hive_bind( :new.message,
                                         hive_binding.type_string,
                                         hive_binding.ref_in );

        bnd.extend;
        bnd( bnd.count ) := hive_bind( to_char( :new.stamp,
                                                  'yyyy-mm-dd' ),
                                         hive_binding.type_date,
                                         hive_binding.ref_in );

    elsif ( deleting ) then

        cmd := q'[ delete from user_log
                    where stamp = ? ]';

        bnd.extend;
        bnd( bnd.count ) := hive_bind( to_char( :new.stamp,
                                                  'yyyy-mm-dd' ),
                                         hive_binding.type_date,
                                         hive_binding.ref_in );

    else

        null; -- should never get here

    end if;

    if ( ( cmd is not null ) and ( bnd.count > 0 ) ) then

        -- execute the remote statement
        hive_remote.dml( cmd, bnd );

    end if;

end user_log_dml;
/

```

This may look more complex, but it's really not. We are simply creating a trigger to handle the DML events and passing them off to Hive-ODCI. The `HIVE_REMOTE.DML()` works in the same way as a query, accepting an array of bind objects and connection information.

The only real difference here is the example is creating and using local variables for `HIVE_BINDS` and the `HIVE_BINDING` types to show how it can be used in that manner.

Hints

The Hive SQL-Like for Hadoop allows for hints to be passed in which control plan execution and join optimization, the Hive-ODCI interface also allows for hints to be passed to the ODCI engine which turn on or provide additional functionality. These hints are parsed and removed from the query before being passed off to Hive for processing. Any hints not supported by Hive-ODCI remain in place and are passed in as provided.

All hints are provided as comment blocks to the query, just like an Oracle or Hive hint. Comment blocks are wrapped, starting with `/*+` and ending in `*/` or start with `--+` and end with a newline.

typecast()

This hint provides a mechanism for casting data types in Hive to ones provided to Oracle. This is useful when the Hive data type is not supported by Oracle or the data type in Hive is insufficient for describing the Oracle type.

Take for example the universal Hive data type `string`, which by default is handled as a `VARCHAR2(4000)` by Hive-ODCI. However, suppose that the original data type was a `CLOB`, or simply that the `string` column exceeds the `4000` character limit? This is where the `typecast()` hint comes into play. You can instruct Hive-ODCI to create a `CLOB` column data type when the column name is encountered with a casting rule.

Casting rules are either space or comma delineated while wrapped in the parenthesis of the `typecast()` keyword. Each rule is formatted by the column name in the Hive record-set delineated by a colon with the following data type and optionally length, precision and scale.

```
column_name:datatype[(length/precision,scale)]
```

Let's take the above SQL and assume that when Scoop was used to originally copied to HDFS that it was done by an over zealous administrator, who took the high-road and made every Hive column a `string`. Oops, that's not going to work for us because the original table had the `STAMP` column as a `DATE` and the `MESSAGE` column as a `CLOB` and we know that it will exceed `4000` characters in length.

So let's confirm by first describing the table ...

```
SQL> col "col_name"    for a30
SQL> col "data_type"   for a30
SQL> col "comment"     for a20

SQL> select * from table( hive_q( 'desc user_log', null, null ) );
```

col_name	data_type	comment
stamp	string	
account	string	
message	string	

If we cast those types to the correct data types, using `typecast()` we can rest assure that the Oracle view or cursor returned is the correct data type.

If we recreate or view, this time using the casting rule, as follows the type defined in Oracle will now be correct.

```
--
create or replace view scott.user_log
(
    stamp,
    account,
    message
)
as
select *
from table( hive_q( q'[ select /*+ stamp:date message:clob */
                        stamp,
                        account,
                        message
                        from user_log
                        order by stamp ]' ) )
/
```

Now when we look at our **VIEW** it is described correctly ...

```
SQL> desc SCOTT.USER_LOG
```

Name	Null?	Type
STAMP		DATE
ACCOUNT		VARCHAR2(30)
MESSAGE		CLOB

For other types, you can optionally define length, precision and scale.
Some other examples may look like ...

```
select /*+
    cust_id:number(9)
    salary:number(7,2)
    dob:timestamp(2)
    first_name:varchar2(50)
    last_name:varchar2(100)
    middle_initial:char(1)
*/
from customers
```

If you have a large number of columns, you can use the helper function **HIVE_HINT()** (below), against a template, to generate the hints for you.

```
SQL> col hint for a80 word_wrap
SQL> select hive_hint( 'SCOTT', 'CUST' ) hint from dual;
```

```
HINT
-----
/*+ cust_id:number(9) salary:number(7,2)
dob:timestamp(2) first_name:varchar2(50)
```

```
last_name:varchar2(100)
middle_initial:char(1) */
```

If the column is not encountered during execution, Hive-ODCI simply ignores the rule.

Final thought

Note we have **not** changed any code in our application or in our PL/SQL procedure. Everything remains exactly as it was before, but our data exists only in Hadoop/Hive. We can view the data, use it with our PL/SQL and even manipulate it with DML. So, you can go to your meeting now and be the hero.

Guidelines

The following sections are guidelines based on practical real-world experience and are intended to help you in the decision making process only. They are not hard-and-fast rules of "thou shalt not" commandments that must be followed to use Hive-ODCI.

The sections are separated into focus areas for Developers and Administrators. In the real-world these are often blurry, completely undefined and are overlapping. While that is true, no matter what role you play, the distinction between them should always be upheld as much as possible.

For Developers

While Hive-ODCI is intended to help the transition of moving data into a Hadoop/Hive data-store, it can with planning and understanding make it transparent to your customers. But at some point it may require Developer intervention.

Performance Considerations

Hive-ODCI is a `PIPELINE` type, so its performance is impacted almost entirely on the performance of Hadoop/Hive. So make sure that Hive is performing at optimal levels.

Ensure that when moving data from Oracle to Hadoop/Hive or creating new tables you have designed in performance from the beginning.

Analytics over in-line views

If you are using in-line views within the query, you may find that using the built-in `RANK()` and `OVER()` functions perform much better.

For example, take the following query

```
select user_log.*
from user_log join
```

```
( select account,
      max( stamp ) as stamp
  from user_log
 group by account ) inline
on user_log.account = inline.account
and user_log.stamp = inline.stamp;
```

While there is nothing technically wrong with this query, based on the plan it could be much better. let's try again ...

```
select *
  from ( select *,
              rank() over
              ( partition by account,
                order by stamp ) as rank
          from user_log ) ranking
 where ranking.rank = 1;
```

Much better. We get the same result with almost a magnitude in increase on performance. If all of this looks familiar, then your right. This has been in Oracle RDBMS for years, so use the same common-sense techniques in your Hadoop/Hive queries as you would in your Oracle RDBMS queries.

The CBO

In recent additions of Hadoop/Hive a Cost-Based-Optimizer (CBO) was introduced, which like in Oracle RDBMS, performs optimizations based on cost, which can adjust the execution plan based on order joins, types of joins, degree of parallelism, etc... leading to increases in query performance.

If not enabled globally, you can still use the Hadoop/Hive CBO, by setting the following parameters at the beginning of the query.

```
set hive.cbo.enable=true;
set hive.compute.query.using.stats=true;
set hive.stats.fetch.column.stats=true;
set hive.stats.fetch.partition.stats=true;
```

Just like Oracle RDBMS, it's important to analyze the tables so the CBO has current cost information. For example, collect statistics at the table and column levels as necessary.

```
analyze table my_table compute statistics;
analyze table my_table compute statistics for columns;
analyze table my_table compute statistics for columns col1, col2;
```

ORCFile

Using ORCFile for Hadoop/Hive table should already be a matter of practice because it is shown to be extremely beneficial in getting fast response times for queries.

If existing tables are not already ORC then it would be prudent to migrate them as soon as possible, even if you convert them by-hand. However, if possible the best case would be to modify the ingest process to use ORC up front.

Apache Tez

Whenever possible use the Apache Tez execution engine instead of a Map-reduce engine. If it is not enabled by default in your environment, then use the following setting at the beginning of the query.

```
set hive.execution.engine=tez;
```

Vectorized queries

Like scans, aggregations, filters and joins vectorized query execution improves performance of operations by splitting them into batches of 1024 rows at a time instead of single row. If not enabled for your environment then use the following setting to at the beginning of the query.

```
set hive.vectorized.execution.enabled = true;  
set hive.vectorized.execution.reduce.enabled = true;
```

New horizons

Putting Hive-ODCI into practice is not all that hard, but it is a new way of doing business. And with "new" things come hiccups and unforeseen circumstances that can put you behind or stop you dead in your tracks.

The biggest thing here, is to identify those risks upfront and mitigate them as soon as possible. Don't wait until the last minute, plan ahead and you won't regret it. I know this is "touchy-feely" advice that you can get at the local coffee shop, but we often forget about the small things in our rush.

Change your PL/SQL code

If your PL/SQL can be changed for the better when using Hive-ODCI, do it. In reviewing your PL/SQL you may find that leveraging the functionality provided by Hive-ODCI makes it more readable, faster, maintainable, etc... then go ahead and make arrangements to modify the existing code base or introduce new code.

Sometimes, that's not feasible or even possible but many times we (the big we) have a tendency to force-fit a solution when that isn't warranted based on the situation.

Keep signatures consistent

Even with the above statement being true, it's best to keep the

underlying `TABLE` an `VIEW` objects with the same signature, column names, types, lengths, etc...

Familiarize yourself with the API

Finally, make sure you familiarize yourself and are comfortable with the Hive-ODCI API. I know reading documentation is boring (writing it even more so, trust me) but take the time and go through it as many times as needed to get a firm understanding. This is particularly important for Developers, as they are the "real" users of the system.

Write some test scripts, see how it behaves in particular situations and make darn sure you know what's coming when you start inserting this stuff into your process.

If you've read the manual, tried it out, found a bug, or simply do not understand, then see the Authors section for POC information and contact me. I'll be happy to help where I can or provide advice and moral support as needed.

For Administrators

Administrators are by enlarge the last line of defense against all manner of issues, they are the gatekeepers for security, performance, storage, new technologies, legacy systems, and a full onslaught of change management from every direction.

Whether the systems are large or small, clustered or single instance the administrator's toolbox is typically a plethora of knowledge, scripts and documentation. Hive-ODCI attempts add to that toolbox, without being just "something else to learn" by leveraging common facilities which already exist for the administrator. Most Hive-ODCI footprint objects will be first-class citizens in the Oracle ecosystem and will be managed, secured and monitored in the same way as other object types.

Wait events

An administrator knows which wait events to look for and which ones can be ignored or are simply part of a running system. Hive-ODCI does not change those presumptions. Mainly because the load Hive-ODCI incurs is on the remote Hadoop/Hive system and not the Oracle RDBMS.

Something to be aware of, however is that unless all your users are hitting Hive-ODCI object simultaneously a wait event for Hive-ODCI will not float to the top of the wait percentages.

A Hive-ODCI object having issue would still be in one of the more innocuous wait events or classes. Additionally, it won't be in a blocking state, so you have to recognize those events which may be indicating problems for the clients.

WAIT_TIME and SECONDS_IN_WAIT

In the `GV$SESSION` view the columns `WAIT_TIME` and `SECONDS_IN_WAIT` can easily be used to determine if Hive-ODCI is waiting on a remote call to respond or complete.

SQL*Net message to client

This wait event can be observed while the Hive-ODCI is sending or responding to the call. While this event itself is not an issue, one that is taking a longer time than normal may indicate that the `query_limit` parameter is set too high.

buffer latch and latch free

While these in normal circumstances, indicate query issues they do not necessarily mean the same thing for Hive-ODCI. You may get these events when the Hive-ODCI is materializing large amounts of data, for example in local sorting or aggregate processing of the returned data.

If this event is observed too high, or too often for a particular Hive-ODCI object then look into off-loading the sort, group or aggregate operations to Hadoop/Hive instead of locally.

Take for example the following

```
select *
  from table( hive_q( q'[ select cust_id,
                           last_name,
                           first_name
                           from cust ]',
                           null,
                           null ) )
 order by cust_id;
```

If this query is showing up in a wait event that is indicating issues materializing the `cust_id` column locally for sorting then considered making the following change to off-load that operation to Hadoop/Hive

```
select *
  from table( hive_q( q'[ select cust_id,
                           last_name,
                           first_name
                           from cust
                           order by cust_id ]',
                           null,
                           null ) );
```

This produces the same end result, but without local materialization and with less wait activity in the Oracle RDBMS.

Storage

Hive-ODCI is a zero storage object solution, like a `VIEW`, for the majority of how its used. However certain aspects of Hive-ODCI will consume space, such as the Hive-ODCI Log and the Saved Filters (Binding).

It is unlikely that your user community, no matter how large, or how much Hive-ODCI is used will generate Filters that impact your storage, necessitating TS extents beyond what was allocated during the installation.

The Hive-ODCI Log data is another story. Depending on the `log_level` set for the System, large amounts of data can be generated. The `log_level` value is a bitmask, which turns types on and off. The types are detailed below but also here

```
create or replace package impl as

--
  none  constant number := 0;
  error constant number := 1;
  warn  constant number := 2;
  info  constant number := 4;
  trace constant number := 8;

  ...

end impl;
```

Each level is progressively verbose with what is written. The types should be obvious, but needless to say `error` writes only critical exceptions, while `trace` writes all operations. So a value of `3` would be `error + warn` and a value of `31` would be `error + warn + info + trace`.

If you find that the Hive Log is filling up faster than expected, review the column `VALUE` in the `DBA_HIVE_PARAMS` to determine the current value.

```
select value
  from dba_hive_params
 where name = 'log_level';
```

If you find that the `log_level` is set appropriately, or as you expected, then this means that the client has set the `session_log_level()` to something too high, which may need to be changed or justified.

The account which created the log data is stored in the column `NAME` found in the `DBA_HIVE_LOG` view. Use the following to get counts of the name generating the most logging information.

```
select name,
```

```

        count(0) total
    from dba_hive_log
    group by name
    order by 2 desc;

```

Log information can be purged by any trusted account granted the `HIVE_ADMIN` role using the following statement.

```
exec dbms_hive.purge_log;
```

Be aware that the statement removes **ALL** log data.

Move to a different TS

By default Hive-ODCI installs to the same tablespace assigned to the `SYSTEM` schema. This for most installations will be sufficient. If you disagree with the installation choice, or you have a larger storage consumption than anticipated you can move Hive-ODCI to another tablespace using the `MOVE_TS` procedure in the `DBMS_HIVE` package.

```

create or replace package dbms_hive as

    ...

    --
    procedure move_ts( ts  in varchar2,
                      obj in varchar2 default null );

end dbms_hive;

```

Purge log and filters

If the Hive-ODCI log data or filter data becomes too large, it can easily be purged using the `PURGE_LOG` and `PURGE_FILTER` procedures respectively in the `DBMS_HIVE` package.

```

create or replace package dbms_hive as

    ...

    --
    procedure purge_log;
    procedure purge_filter( key in varchar2 default null );

    ...

end dbms_hive;

```

These procedures do exactly what they indicate. And because they are exposed through the Hive-ODCI management package they can be called by a trusted user who has been granted the `HIVE_ADMIN`

role.

Role Based Access Control

As indicated multiple times now, Hive-ODCI objects are primarily first-class and are managed with RBAC just like any other object.

When Hive-ODCI is installed, it creates 2 helper roles for you setting up the appropriate permissions for their use.

The `HIVE_USER` role has access to the common Hive-ODCI objects making it possible to Query, use DML or DDL, manage Bind variables, etc...

The second role, is `HIVE_ADMIN` which provides access to the management of Hive-ODCI itself, such as Logs, Parameters, etc... This role is also granted to the `DBA` role by default so be cognizant that a user with `HIVE_ADMIN` is considered a trusted and responsible party, similar to a user with `DBA` but having less permissions in the database.

Bindings

Filters, or saved `HIVE_BINDS` work more like Network ACL (see `DBMS_NETWORK_ACL_ADMIN`) and Java Policies (see `DBMS_JAVA`) than they do first-class object permissions.

When created a Filter will by default assign ownership of the `HIVE_BINDS` object key to the account which created it.

That account, or a `HIVE_ADMIN` assigned account can grant and revoke permissions to both users and roles using the `HIVE_BINDING` type methods.

```
create or replace package binding as

...

--
priv_read      constant guard      := 1;
priv_write     constant guard      := 2;
priv_readwrite constant guard      := 3;

--
procedure allow( key in varchar2,
                 act in varchar2,
                 lvl in guard default priv_readwrite );

--
procedure deny( key in varchar2,
                act in varchar2 );

...

end binding;
```

The `ALLOW()` procedure works like a `GRANT` and the `DENY()` works like a `REVOKE`

When a `HIVE_BIND` is accessed from the stored Filter the permission is checked for the operation (e.g. read, write, etc...). If a `PUBLIC` grant is made, then all accounts may access the Filter at the level granted.

Java Policies

Hive-ODCI will in no way automatically grant or revoke Java Policy permissions, therefore it is necessary for the administrator to take on this action as needed.

When Hive-ODCI attempts to set a Java Property at run-time, which has not been granted the appropriate policy an exception is thrown. As part of that exception, the necessary `DBMS_JAVA.GRANT_PERMISSION()` is provided to the caller as part of the exception text. It is likely that the text generated by the database will be sent to the administrator requesting it be executed.

Before doing so, make sure there is a clear understanding of what that permission provides as well as what the property enables or achieves.

Failure to understand the property and/or permission may result in reducing the security posture of the database or at worst make the database unstable. For the most part this will not happen, but it is good practice to have full situational awareness before proceeding.

API

The Application programming Interface (API) for Hive-ODCI is accessible through the PL/SQL objects created during installation. The objects include Views, Packages, Procedure, Types and Functions each providing a unique set of functionality based on need.

Views

Views are the insight into Hive-ODCI, providing invaluable information for both Administrators and Users.

dba_hive_params

Lists both the system and session level values.

This view is accessible only by the `HIVE_ADMIN` role

- Columns

<code>name</code>	- Parameter name
<code>session_value</code>	- Current session value
<code>system_value</code>	- Default system value

dba_hive_filters

Lists the saved Filters.

This view is accessible only by the `HIVE_ADMIN` role

- Columns

key	-	Filter key name
seq	-	The ordinal <code>bind</code> sequence
type	-	The <code>bind type</code>
scope	-	The <code>bind</code> reference scope
value	-	The value of the <code>bind</code> data

dba_hive_filter_privs

Lists the privileges of the saved Filters.

This view is accessible only by the `HIVE_ADMIN` role

- Columns

key	-	Filter key name
grantee	-	The grantee name, user or role
read	-	Contains read access
write	-	Contains write access

dba_hive_log

Lists the log data

This view is accessible only by the `HIVE_ADMIN` role

- Columns

stamp	-	The time stamp the log record
name	-	Account name who create the log record
type	-	The type of log record
tier	-	The <code>text</code> representation of the type
<code>text</code>	-	The <code>text</code> of the log record

user_hive_log

Lists the log data for the current user

This view is accessible through the `HIVE_USER` role

- Columns

stamp	-	The time stamp the log record
type	-	The type of log record
tier	-	The text representation of the type
text	-	The text of the log record

user_hive_params

Lists current parameter values used by the session

This view is accessible by both the `HIVE_ADMIN` and `HIVE_USER` roles

- Columns

<code>name</code>	-	Parameter name
<code>value</code>	-	Parameter <code>value</code>

user_hive_filters

Lists the saved Filters which are accessible by the session

This view is accessible by both the `HIVE_ADMIN` and `HIVE_USER` roles

- Columns

<code>key</code>	-	Filter key name
<code>seq</code>	-	The ordinal <code>bind</code> sequence
<code>type</code>	-	The <code>bind type</code>
<code>scope</code>	-	The <code>bind</code> reference scope
<code>value</code>	-	The value of the <code>bind</code> data

user_hive_filter_privs

Lists the privileges of the saved Filters which are accessible by the session

This view is accessible by both the `HIVE_ADMIN` and `HIVE_USER` roles

- Columns

<code>key</code>	-	Filter key name
<code>grantee</code>	-	The grantee name, user or role
<code>read</code>	-	Contains <code>read</code> access
<code>write</code>	-	Contains <code>write</code> access

Packages

binding

An interface for creating, manipulating and saving `HIVE_BINDS` arrays of `HIVE_BIND` objects. These objects are used in the marshaling and interpretation of arguments to and from the JDBC Driver providing both data type and direction.

This object is accessible through the `HIVE_USER` role and is

referenced as the `SYNONYM` name `HIVE_BINDING`

Subtypes

- reference -- `number` for bind scope values
- typeof -- `number` for bind value types
- guard -- `number` for bind access privileges

Constants

Helper constants for abstracting type information

Generic

Name	Value

<code>none</code>	<code>0</code>

Privilege

Name	Value

<code>priv_read</code>	<code>1</code>
<code>priv_write</code>	<code>2</code>
<code>priv_readwrite</code>	<code>3</code>

Scope

Name	Value

<code>scope_in</code>	<code>1</code>
<code>scope_out</code>	<code>2</code>
<code>scope_inout</code>	<code>3</code>

Type

Name	Value

<code>type_bool</code>	<code>1</code>
<code>type_date</code>	<code>2</code>
<code>type_float</code>	<code>3</code>
<code>type_int</code>	<code>4</code>
<code>type_long</code>	<code>5</code>
<code>type_null</code>	<code>6</code>
<code>type_rowid</code>	<code>7</code>
<code>type_short</code>	<code>8</code>
<code>type_string</code>	<code>9</code>
<code>type_time</code>	<code>10</code>
<code>type_timestamp</code>	<code>11</code>
<code>type_url</code>	<code>12</code>

Interface

get()

Overloaded `FUNCTION` to retrieve a saved `HIVE_BINDS`

filter based on Key name or returns a single `HIVE_BIND` from the session array.

- Prototype

```
function get( key in varchar2 ) return binds;

function get( idx in number, lst in binds ) return bind;
```

- Parameter

key	-	Saved key name
idx	-	Ordinal number of the item.

count()

Overloaded `FUNCTION` to count the length of a saved `HIVE_BINDS` filter based on Key name or an existing one provided.

- Prototype

```
function count( key in varchar2 ) return number;
function count( lst in binds ) return number;
```

- Parameter

key	-	Saved key name
lst	-	Existing <code>HIVE_BINDS</code> array

new()

A `FUNCTION` to create a new, single, `HIVE_BIND` object.

- Prototype

```
function new( value in varchar2,
              type in typeof default type_string,
              scope in reference default scope_in ) return bind;
```

- Parameter

value	-	Value of the <code>HIVE_BIND</code>
type	-	Interpretation type of the <code>HIVE_BIND</code>
scope	-	Reference scope, direction, for using the <code>HIVE_BIND</code>

append()

Overloaded `PROCEDURE` to append a `HIVE_BIND` object to an existing `HIVE_BINDS` array, saved or passed in.

- Prototype

```

procedure append( key    in varchar2,
                  value in varchar2,
                  type  in typeof    default type_string,
                  scope in reference default scope_in );

procedure append( value in    varchar2,
                  type  in    typeof    default type_string,
                  scope in    reference default scope_in,
                  lst   in out binds );

procedure append( key in varchar2, val in bind );

procedure append( val in bind, lst in out binds );

procedure append( key in varchar2, val in binds );

procedure append( val in binds, lst in out binds );

```

- Parameter

key	-	Saved key name
value	-	Value of the HIVE_BIND
type	-	Interpretation type of the HIVE_BIND
scope	-	Reference scope, direction, for using the HIVE_BIND
lst	-	Existing HIVE_BINDS array
val	-	A HIVE_BINDS array to be appended to the existing array

change()

A **PROCEDURE** to modify an exiting element of a saved **HIVE_BIND** object.

- Prototype

```

procedure change( key    in varchar2,
                  idx   in number,
                  value in varchar2,
                  type  in typeof    default type_string,
                  scope in reference default scope_in );

```

- Parameter

key	-	Saved key name of the array
idx	-	Ordinal number of the item to modify
value	-	Value of the HIVE_BIND
type	-	Interpretation type of the HIVE_BIND
scope	-	Reference scope, direction, for using the HIVE_BIND

remove()

A **PROCEDURE** to delete an exiting element of a saved **HIVE_BIND** object.

- Prototype

```
procedure remove( key in varchar2,  
                  idx in number );
```

- Parameter

key	-	Saved key name of the array
idx	-	Ordinal number of the item to modify

replace()

A **PROCEDURE** to replace one **HIVE_BINDS** array with another.

- Prototype

```
procedure replace( val in binds, lst in out binds );
```

- Parameter

val	-	A HIVE_BINDS array to be appended to the existing array
lst	-	HIVE_BINDS array to overwrite

clear()

An overloaded **PROCEDURE** to remove all elements in a saved Filter or a passed in array.

- Prototype

```
procedure clear( key in varchar2 );  
procedure clear( lst in out binds );
```

- Parameter

key	-	Saved key name of the array
lst	-	HIVE_BINDS array to overwrite

allow()

A **PROCEDURE** to grant access rights to saved Filter. If an existing grant was already made, then it is replaced by the change

- Prototype

```
procedure allow( key in varchar2,  
                act in varchar2,  
                lvl in guard default priv_readwrite );
```

- Parameter

key	-	Saved key name of the array
-----	---	-----------------------------

- act - Name **of** the account to allow access, which can be an Oracle username **or** role
- lvl - The level **of** access to grant

deny()

A **PROCEDURE** to revoke access rights from a saved Filter. If no existing grant exists an exception is thrown.

- Prototype

```
procedure deny( key in varchar2,  
               act in varchar2 );
```

- Parameter

- key - Saved key name **of** the array
- act - Name **of** the account to allow access, which can be an Oracle username **or** role

save()

A `PROCEDURE` to save a passed in `HIVE_BINDS` array. If a key of the same name already exists, an exception is thrown.

- Prototype

```
procedure save( key in varchar2, lst in binds );
```

- Parameter

```
key      -   Saved key name of the array
lst      -   HIVE_BINDS array to save
```

Exception

ex unknown

This exception is thrown when an unknown error has been encountered.

```
ex_unknown    exception;  
es_unknown    constant varchar2( 256 ) := 'Unknown error encountered';  
ec_unknown    constant number := -20001;  
pragma        exception_init( ex_unknown, -20001 );
```

ex_denied

This exception is thrown when a saved Filter access attempt has failed because of insufficient privileges.

```
ex_denied    exception;
es_denied    constant varchar2( 256 ) := 'Request denied,  
insufficient privileges';
```

```
ec_denied    constant number := -20002;
pragma      exception_init( ex_denied, -20002 );
```

ex_no_grant

This exception is thrown when a revoke attempt is made on a saved Filter that does not contain a previous grant for the account specified

```
ex_no_grant exception;
es_no_grant constant varchar2( 256 ) := 'Privileges not granted';
ec_no_grant constant number := -20003;
pragma      exception_init( ex_no_grant, -20003 );
```

dbms_hive

A `PACKAGE` providing the interfaces for management of Hive-ODCI configurations.

This object is accessible through the `HIVE_ADMIN` role and is referenced as the `SYNONYM` name `DBMS_HIVE`

Interface

exist()

A `FUNCTION` returning a `BOOLEAN` value
(`true` or `false`) if the parameter name exists

- Prototype

```
function exist( name in varchar2 ) return boolean;
```

- Parameter

```
name      - Case sensitive name of the parameter to check
```

param()

An overloaded `FUNCTION` and `PROCEDURE` to both set and get a parameter. When setting, if the parameter exists then it's value is overwritten

- Prototype

```
-- get
function param( name in varchar2 ) return varchar2;

-- set
procedure param( name in varchar2, value in varchar2 );
```

- Parameter

```
name      - Case sensitive name of the parameter to get
```

value - or **set**
Value of the parameter **when** setting. The **value**
is the **return** for the **FUNCTION**

If the a parameter does not exist, then the function returns `NULL`

remove()

A `PROCEDURE` for removing a parameter. If the parameter does not exist, no error is thrown.

- Prototype

```
procedure remove( name in varchar2 );
```

- Parameter

name - Case sensitive name **of** the parameter to
delete

purge_log()

A `PROCEDURE` which purges all data in the `DBA_HIVE_LOG` optionally removing single user records. The `name` parameter is case-sensitive and when provided removes records matching only the those records created by `name` . If no `name` is provided or `NULL` is provided the table is truncated and storage dropped (`TRUNCATE TABLE LOG$ DROP STORAGE`). This is to provide a faster cleanup procedure than issuing an unqualified `DELETE` statement.

- Prototype

```
procedure purge_log( name in varchar2 default null );
```

purge_filter()

A `PROCEDURE` for removing all saved Filters by name. If no name value or a `NULL` is provided as the argument then **ALL** Filters are removed.

- Prototype

```
procedure purge_filter( key in varchar2 default null );
```

- Parameter

key - Key name **of** the array to remove, **this** defaults to `NULL` which indicates ALL values are removed

move_ts()

A `PROCEDURE` moving the Hive-ODCI objects to different

tablespace. If the current tablespace name is the same as the one specified then no errors are thrown. If the tablespace specified does not exist then an exception is thrown.

An `OBJECT` name is alternatively provided which will move only a single object. If no value or `NULL` is used then **ALL** objects are moved.

Only the values "`NULL`", "`param`", "`filter`", "`priv`" or "`log`" can be used, all other values will be ignored.

- Prototype

```
procedure move_ts( ts   in varchar2,
                  obj in varchar2 default null );
```

- Parameter

```
ts      - Destination tablespace name.
obj     - Optional object name, can be NULL, "param", "filter",
          "priv" or "log". Any other value will be ignored
```

remote

This `PACKAGE` is the interface for managing the remote connectivity of the current session. This is also the interface to execute remote commands, such as queries, DDL and DML operations.

This object is accessible through the `HIVE_USER` role and is referenced as the `SYNONYM` name `HIVE_REMOTE`

Interface

`session_param()`

Overloaded `FUNCTION` and `PROCEDURE` to get and set parameters only for the current session. Parameters set using this `PROCEDURE` are valid only for the duration of the Oracle session and are discarded when the session ends.

A session level parameter takes precedence over the system level parameter of the same name.

- Prototype

```
-- get
function session_param( name in varchar2 ) return varchar2;

-- set
procedure session_param( name in varchar2,
                        value in varchar2 );
```

- Parameter

name	-	Case sensitive name of the parameter to get or set
value	-	Value of the parameter when setting. The value is the return for the FUNCTION

If the session level parameter is not set, or does not exist then the function returns `NULL`

session_log_level()

A `PROCEDURE` which allows the logging level to be set for the current session. This setting is valid only for the duration of the Oracle session and are discarded when the session ends.

- Prototype

```
procedure session_log_level( typ in number );
```

- Parameter

typ	-	The log level type
------------	---	--------------------

Log level values are bit masks, to allow for multiple values to be specified in a single numeric value. They are defined as part of the `IMPL` package specification.

```
none constant number := 0;
error constant number := 1;
warn constant number := 2;
info constant number := 4;
trace constant number := 8;
```

session_clear()

This `PROCEDURE` clears the session `CONNECTION` type.

- Prototype

```
procedure session_clear;
```

After being cleared, consecutive attempts will rebuild the `CONNECTION` type from the parameter information.

session()

An overloaded `FUNCTION` and `PROCEDURE` which creates, gets, or modifies the current session `CONNECTION` type.

- Prototype

```
-- get
function session return connection;
```

```
-- set
procedure session( url in varchar2 );

procedure session( usr in varchar2,
                  pwd in varchar2 );

procedure session( url in varchar2,
                  usr in varchar2,
                  pwd in varchar2 );

procedure session( url in varchar2,
                  usr in varchar2,
                  pwd in varchar2,
                  ath in varchar2 );

procedure session( con in connection );
```

- Parameter

url	-	The JDBC URL of the remote connection
usr	-	User name to be applied for the connection
pwd	-	The user password to be applied
ath	-	The authentication type of the connection

When getting the current `CONNECTION` type, the password value is redacted by the function, no matter what its value or the authentication type specified.

query()

A `PIPELINED FUNCTION` which returns an `ANYDATASET` value. This value contains the description information of a given row type along with the set of data instances of that row.

- Prototype

```
function query( stm in varchar2,
               bnd in binds      default null,
               con in connection default null )
return anydataset pipelined using hive_t;
```

- Parameter

stm	-	The SQL query statement to be executed remotely
bnd	-	Optionally, the <code>HIVE_BINDS</code> array containing the values to be bound to the SQL statement
con	-	Optionally, the remote <code>CONNECTION</code> information

This function can be used in a casting statement to retrieve the remote records, but it does not contain a `PACKAGE BODY` definition, rather it uses the `HIVE_T` ODCI object type making the following 2 examples equivalent

```
val := hive_q( 'select cust_id,
```



```
        last_name,  
        first_name  
    from cust', null, null );
```

vs.

```
val := hive_remote.query( 'select cust_id,  
                          last_name,  
                          first_name  
                          from cust', null, null );
```

dml()

This `PROCEDURE` allows remote DML to be executed using the provided SQL statement, bind variables and connection type.

- Prototype

```
procedure dml( stm in varchar2,  
              bnd in binds      default null,  
              con in connection default null );
```

- Parameter

stm	-	The DML statement to be executed remotely
bnd	-	Optionally, the HIVE_BINDS array containing the values to be bound to the DML statement
con	-	Optionally, the remote CONNECTION information

ddl()

This `PROCEDURE` allows remote DDL to be executed using the provided command statement and connection type.

- Prototype

```
procedure ddl( stm in varchar2,  
              con in connection default null );
```

- Parameter

stm	-	The DML statement to be executed remotely
con	-	Optionally, the remote CONNECTION information

impl

This `PACKAGE` is the primary implementation of the ODCI layer and the interface called by Oracle ODCI through the `HIVE_T` object type.

This interface is used internally only, and while it is available for execution with the `HIVE_ADMIN` role it should never be called

directly. It is included here only for completeness of the documentation.

There is no SYNONYM for this object and both its specification and body are wrapped during the installation process.

Constants

Helper constants for abstracting type information

Log level

```
none constant number := 0;  
error constant number := 1;  
warn constant number := 2;  
info constant number := 4;  
trace constant number := 8;
```

Log level values are bit masks, to allow for multiple values to be specified in a single numeric value.

Interface

log()

This procedure writes a logging record of the type specified, if `log_level` at the systems or session set set to support it.

- Prototype

```
procedure log( typ in number, txt in varchar2 );
```

Additionally, there are pass-through procedures which call this procedure using the predefined type.

```
procedure log_error( txt in varchar2 );  
procedure log_warn( txt in varchar2 );  
procedure log_info( txt in varchar2 );  
procedure log_trace( txt in varchar2 );
```

- Parameter

typ	-	The type of log to write, ignored if the log_level is not set to support that type
txt	-	The text of the message

session_param()

Overloaded `FUNCTION` and `PROCEDURE` to get and set parameters only for the current session. Parameters set using this `PROCEDURE` are valid only for the duration of the Oracle session and are discarded when the session ends.

A session level parameter takes precedence over the system level parameter of the same name.

- Prototype

```
-- get
function session_param( name in varchar2 ) return varchar2;

-- set
procedure session_param( name in varchar2,
                        value in varchar2 );
```

- Parameter

name	-	Case sensitive name of the parameter to get or set
value	-	Value of the parameter when setting. The value is the return for the FUNCTION

If the session level parameter is not set, or does not exist then the function returns `NULL`

session_log_level()

A `PROCEDURE` which allows the logging level to be set for the current session. This setting is valid only for the duration of the Oracle session and are discarded when the session ends.

- Prototype

```
procedure session_log_level( typ in number );
```

- Parameter

typ	-	The log level type
-----	---	--------------------

session_clear()

This `PROCEDURE` clears the session `CONNECTION` type.

- Prototype

```
procedure session_clear;
```

After being cleared, consecutive attempts will rebuild the `CONNECTION` type from the parameter information.

session()

An overloaded `FUNCTION` and `PROCEDURE` which creates, gets, or modifies the current session `CONNECTION` type.

- Prototype

```

-- get
function session return connection;

-- set
procedure session( url in varchar2 );

procedure session( usr in varchar2,
                   pwd in varchar2 );

procedure session( url in varchar2,
                   usr in varchar2,
                   pwd in varchar2 );

procedure session( url in varchar2,
                   usr in varchar2,
                   pwd in varchar2,
                   ath in varchar2 );

procedure session( con in connection );

```

- Parameter

url	-	The JDBC URL of the remote connection
usr	-	User name to be applied for the connection
pwd	-	The user password to be applied
ath	-	The authentication type of the connection

When getting the current `CONNECTION` type, the password value is **not** redacted by this function, unlike the function of the same name used in the `REMOTE` package

sql_describe()

This overloaded `FUNCTION` is called by ODCI to describe information for a table whose return type is `ANYDATASET`. This function is called when `HIVE_T` `ODCITableDescribe` is encountered.

- Prototype

```

function sql_describe( stm in varchar2,
                      bnd in binds      default null,
                      con in connection default null ) return anytype;

function sql_describe( typ out anytype,
                      stm in  varchar2,
                      bnd in  binds      default null,
                      con in  connection default null ) return number;

function sql_describe( key in  number,
                      typ out anytype ) return number;

```

- Parameter

stm	-	The SQL query statement to be executed remotely
bnd	-	Optionally, the HIVE_BINDS array containing the

	values to be bound to the SQL statement
con	- Optionally, the remote CONNECTION information
typ	- The ANYTYPE value that describes the returned rows from the table function
key	- The transient key value used in the concurrent calling chain, identifying the record of the context

If successful the return value for each function will be

`ODCICONST.SUCCESS` otherwise it will return `ODCICONST.ERROR`

sql_open()

This `FUNCTION` is called by ODCI to initialize the scan of a table function. This function is called when `HIVE_T` `ODCITableStart` is encountered.

- Prototype

```
function sql_open( key out number,
                  stm in  varchar2,
                  bnd in  binds      default null,
                  con in  connection default null ) return number;
```

- Parameter

key	- The transient key value used in the concurrent calling chain, identifying the record of the context
stm	- The SQL query statement to be executed remotely
bnd	- Optionally, the HIVE_BINDS array containing the values to be bound to the SQL statement
con	- Optionally, the remote CONNECTION information

If successful the return value for each function will be

`ODCICONST.SUCCESS` otherwise it will return `ODCICONST.ERROR`

sql_fetch()

This `FUNCTION` is called by ODCI to next batch of rows from a table function. This function is called when `HIVE_T` `ODCITableFetch` is encountered.

- Prototype

```
function sql_fetch( key in  number,
                   num in  number,
                   rws out records ) return number;
```

- Parameter

key	- The transient key value used in the concurrent calling chain, identifying the record of the context
num	- The number of rows the system expects in the current fetch cycle.
rws	- The RECORDS array for the rows requested in the

current fetch cycle.

If successful the return value for each function will be

`ODCICONST.SUCCESS` otherwise it will return `ODCICONST.ERROR`

sql_close()

This `FUNCTION` is called by ODCI to performs cleanup operations after scanning a table function cycle is complete. This function is called when `HIVE_T` `ODCITableClose` is encountered.

- Prototype

```
function sql_close( key in number ) return number;
```

- Parameter

key	-	The transient key value used in the concurrent calling chain, identifying the record of the context
-----	---	---

If successful the return value for each function will be

`ODCICONST.SUCCESS` otherwise it will return `ODCICONST.ERROR`

sql_dml()

This `PROCEDURE` allows remote DML to be executed using the provided SQL statement, bind variables and connection type.

- Prototype

```
procedure sql_dml( stm in varchar2,  
                  bnd in binds      default null,  
                  con in connection default null );
```

- Parameter

stm	-	The DML statement to be executed remotely
bnd	-	Optionally, the HIVE_BINDS array containing the values to be bound to the DML statement
con	-	Optionally, the remote CONNECTION information

sql_ddl()

This `PROCEDURE` allows remote DDL to be executed using the provided command statement and connection type.

- Prototype

```
procedure sql_ddl( stm in varchar2,  
                  con in connection default null );
```

- Parameter

```
stm      - The DML statement to be executed remotely
con      - Optionally, the remote CONNECTION information
```

Exception

ex_not_eligible

This exception is thrown when a request to change a parameter at the session level is ineligible for the operation.

```
ex_not_eligible exception;
es_not_eligible constant varchar2( 256 ) := 'Parameter is not
                                             eligible for
                                             change at the
                                             session level';
ec_not_eligible constant number           := -20103;
```

Functions

Only the `HIVE_Q` function is available for execution, while all other functions are for internal use only.

hive_q

A `PIPELINED FUNCTION` which returns an `ANYDATASET` value. This value contains the description information of a given row type along with the set of data instances of that row.

This function is accessible through the `HIVE_USER` role and is referenced as the `SYNONYM` name `HIVE_Q`

- Prototype

```
function hive_q( stm in varchar2,
                 bnd in binds      default null,
                 con in connection default null )
return anydataset pipelined using hive_t;
```

- Parameter

```
stm      - The SQL query statement to be executed remotely
bnd      - Optionally, the HIVE_BINDS array containing the
           values to be bound to the SQL statement
con      - Optionally, the remote CONNECTION information
```

This function can be used in a casting statement to retrieve the remote records, but it does not contain a `PACKAGE BODY` definition, rather it uses the `HIVE_T` ODCI object type making the following 2 examples equivalent

```
val := hive_q( 'select cust_id,
               last_name,
```

```
        first_name
    from cust', null, null );
```

vs.

```
val := hive_remote.query( 'select cust_id,
                           last_name,
                           first_name
                           from cust', null, null );
```

hive_hint

This helper `FUNCTION`, accessible through the `HIVE_USER` role, returns a hints for an `OWNER.TABLE` column listing. The hint string returned can be passed into a select for a `HIVE_Q` query. Please see *Hints* section above for more information

- Prototype

```
function hive_hint( own in varchar2,
                   tab in varchar2,
                   typ in varchar2 default 'typecast' )
```

- Parameter

own	-	Owner of the template table to generate hints from
tab	-	Table template to generate hints from
typ	-	Type of hint string to generate

bitxor

This `FUNCTION` performs the logical exclusive OR operation on each pair of corresponding bits provided.

- Prototype

```
function bitxor( x in number, y in number ) return number;
```

- Parameter

x	-	First bit mask
y	-	Second bit mask

bitnot

This `FUNCTION` performs logical negation on each bit, forming the complement of the given binary value provided.

- Prototype

```
function bitnot( x in number ) return number;
```


- Parameter

x - Bit mask

oid

This **FUNCTION** returns the database object identifier for the name provided, which is expected to be a **USERNAME** or **ROLE**

The return value is used as the join key for the saved Filter privileges

- Prototype

```
function oid( o in varchar2 ) return number;
```

- Parameter

o - Object name

oname

This **FUNCTION** returns the database object name for the identifier provided, which is expected to be a **USER#** representing the **USERNAME** or **ROLE**

The return value is used as the join key displaying names for the saved Filter privileges.

- Prototype

```
function oname( o in number ) return varchar2;
```

- Parameter

o - Object identifier

Types

Hive-ODCI types are both objects and table (arrays) and are used in the transparent conversions from the JDBC Driver to the PL/SQL calls.

attribute

This **OBJECT** type is used to describe a column value. it is used internally only to define rows in a record set and a table description.

- Members

name	- Name of the column
code	- Data type code of the column
prec	- The precision of the column
scale	- The scale of the column
len	- The length of the column
csid	- The locale identifier
csfrm	- The locale format

attributes

A `TABLE` array of `ATTRIBUTE` objects.

- Members

table **of** attribute

data

This `OBJECT` type contains a single column of information.

- Members

code	- The data type code, pointing to the member which actually contains the data
val_varchar2	- String data
val_number	- Numeric data
val_date	- Date data
val_timestamp	- Timestamp data
val_clob	- CLOB data
val_blob	- BLOB data

records

A `TABLE` array of `DATA` objects, representing a single row

- Members

table **of** data

connection

This `OBJECT` type contains the remote connection information

- Members

url	- The JDBC URL of the remote connection
name	- User name to be applied for the connection
pass	- The user password to be applied
auth	- The authentication type of the connection

bind

This `OBJECT` defines the bind variable used for the remote command

- Members

value	- The value of the data to be bound
type	- The data type code, for how to bind the data
scope	- The reference code direction of the bind

binds

A **TABLE** array of **BIND** objects

- Members

table **of** bind

hive_t

An **OBJECT** type for the data cartridge abstraction later. This object is the compliant interface for ODCI callback functionality.

- Members

key	- The transient key value used in the concurrent calling chain, identifying the record of the context
ref	- The persistent ANYTYPE reference being populated in the current calling cycle

ODCITableDescribe

Retrieves describe information for a table function whose return type is **ANYDATASET** .

- Prototype

```
function ODCITableDescribe( typ out anytype,
                           stm in varchar2,
                           bnd in binds      := null,
                           con in connection := null ) return number
```

- Parameter

typ	- The AnyType value that describes the returned rows from the table function
stm	- The command statement to be executed remotely
bnd	- Optionally, the HIVE_BINDS array containing the values to be bound to the statement
con	- Optionally, the remote CONNECTION information

ODCITablePrepare

This **FUNCTION** prepares the scan context and command information upon a compile request.

- Prototype

```
function ODCITablePrepare( ctx out hive_t,  
                           inf in sys.ODCITabFuncInfo,  
                           stm in varchar2,  
                           bnd in binds      := null,  
                           con in connection := null ) return number
```

- Parameter

ctx	-	The scan context created by this routine is the value passed in as a parameter to the later routines in the command cycle
inf	-	The projection information and the return descriptor object
stm	-	The command statement to be executed remotely
bnd	-	Optionally, the HIVE_BINDS array containing the values to be bound to the statement
con	-	Optionally, the remote CONNECTION information

ODCITableStart

This **FUNCTION** initializes the scan of a table function to start the command cycle.

- Prototype

```
function ODCITableStart( ctx in out hive_t,  
                         stm in      varchar2,  
                         bnd in      binds      := null,  
                         con in      connection := null ) return number
```

- Parameter

ctx	-	The scan context modified by this routine is the value passed in as a parameter to the later routines in the command cycle
stm	-	The command statement to be executed remotely
bnd	-	Optionally, the HIVE_BINDS array containing the values to be bound to the statement
con	-	Optionally, the remote CONNECTION information

ODCITableFetch

This **FUNCTION** returns the next batch of rows in the command cycle.

- Prototype

```
function ODCITableFetch( cts in out hive_t,  
                         num in      number,  
                         rws out      anydataset ) return number
```

- Parameter

ctx	-	The scan context modified by this routine is the value passed in as a parameter to the later routines in the command cycle
num	-	The number of rows the system expects in the current fetch cycle.
rws	-	The RECORDS array for the rows requested in the current fetch cycle.

ODCITableClose

This `FUNCTION` performs cleanup operations after command cycle is complete

- Prototype

```
function ODCITableClose( ctx in hive_t ) return number
```

- Parameter

ctx	-	The scan context created by this routine is the value passed in as a parameter to the later routines in the command cycle
-----	---	--

Parameters

Parameters provide the customization of Hive-ODCI.

application

The application name

version

The installed or patched version

license

The BSD licensing agreement

log_level

The current log level

hive_jdbc_driver

The fully qualified Java class name of the driver to load

hive_jdbc_url

The JDBC URL for the remote system

hive_jdbc_url.X

These parameters are additional URL key value pairs in consecutive order. A gap in the numbering sequence will cause Hive-ODCI to stop reading the parameters assuming that it has reached the end of the list.

java_property.X

These parameters are the Java System properties to set in consecutive order. A gap in the numbering sequence will cause Hive-ODCI to stop reading the parameters assuming that it has reached the end of the list.

hive_user

The `Driver.getConnection()` call in Java is overloaded to accept a URL only or optionally with a User/Password. When set this parameter is passed through to the call as the user value.

hive_pass

The `Driver.getConnection()` call in Java is overloaded to accept a URL only or optionally with a User/Password. When set this parameter is passed through to the call as the password value.

hive_auth

The authentication type for the JDBC Driver.

query_limit

The ceiling limitation for any given query. When set no query will return more rows than the parameter specified

Roles

hive_user

This role provides the access privileges necessary for common Hive-ODCI functionality.

hive_admin

This role provides the escalated access privileges necessary used by administrators of Hive-ODCI.

Synonyms

Hive-ODCI synonyms are the public alternative names for the interface objects providing location transparency.

- hive_q
- hive_t
- hive_remote
- hive_bind
- hive_binds

- hive_binding
- hive_attribute
- hive_attributes
- hive_data
- hive_records
- hive_connection
- dbms_hive
- dba_hive_params
- dba_hive_filters
- dba_hive_filter_privs
- dba_hive_log
- user_hive_params
- user_hive_filters
- user_hive_filter_privs