

# Hive-ODCI - Users Guide

---

Hive-ODCI is an [Oracle Data Cartridge Interface](#) for dynamically accessing Hadoop/Hive data-stores through an Oracle 12c database. In other words Hive-ODCI makes Hadoop/Hive tables accessible as first-class, native, objects directly using PL/SQL, SQL, VIEWS, DML, DDL, etc.... in an Oracle 12c database.

---

## Author

---

Metasystem Technologies Inc. (MTI)  
[www.mtihn.com](http://www.mtihn.com)

Nicholas Van Wyen  
nvanwyen@mtihn.com

## License

---

**Hive-ODCI - Copyright (C) 2006-2016 Metasystems Technologies Inc. (MTI)  
Nicholas Van Wyen**

This library is free software; you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation; either version 2.1 of the License, or (at your option) any later version.

This library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this library; if not, write to the

Free Software Foundation, Inc.  
59 Temple Place, Suite 330,  
Boston, MA 02111-1307 USA

## Releases

---

All releases can be found on Github  
<https://github.com/nvanwyen/hive-odci/releases>, along with the [latest release](#).

The project home is publicly available on Github at  
<https://github.com/nvanwyen/hive-odci>

# Installation and Removal

---

See `INSTALL.md` for instructions

## Concepts

---

Hive-ODCI is a pass-through interface allowing SQL access from within an Oracle RDBMS to information retained in an external Hive/Hadoop data-store. Hive-ODCI provides PL/SQL interfaces using ODCI to accomplish this functionality, making the access viable as native object in Oracle.

Hive-ODCI is accessed via the `HIVE` schema, and controlled through RBAC permissions in Oracle. A client (user or application) is granted privileges through the `HIVE_USER` role or direct system privileges by the DBA.

The client uses the PL/SQL objects to query, or execute DML/DDDL in the remote Hive datastore, using the `HIVE_T` object type or one of the PL/SQL packages.

The `HIVE` schema installs Java classes in the database, which perform the JDBC execution on the clients behalf, and streams `PIPLINED` results back through the same interface.

The client can dictate most levels of functionality at run-time, with predefined session data, bind-variables, etc... as customization for each call.

Take for example the following concepts. The Hive/Hadoop data-store is remotely accessible on a separate server. The Hive-ODCI Java classes access the remote data via the JDBC Driver loaded during installation.

The client accesses the Hive-ODCI interface using the PL/SQL objects provided and/or a first-class `VIEW`, controlled by RBAC, for a 2-way avenue for data.



Our tables looks like this ...

```
SQL> desc SCOTT.USER_LOG
```

| Name    | Null?    | Type           |
|---------|----------|----------------|
| STAMP   | NOT NULL | DATE           |
| ACCOUNT |          | VARCHAR2(30)   |
| MESSAGE |          | VARCHAR2(4000) |

Because we have tons of room available in our Hadoop cluster we decide that we want to move the data there, so it can be indexed and searched. But hold on, we have a problem, we still have an application that reads the table and creates reports for upper management and they are not going to change their application to read from 2 different places using 2 different methods (oh, what to do).

## Hive-ODCI to the rescue

On top of inserting, updating and deleting capabilities to the table, the application also contains PL/SQL to create the reports and has a VIEW used in displaying the monthly metrics.

They look something like this ...

```
--
procedure user_log_report( p_report out xmltype ) is
begin

    for rec in ( select account,
                      message
                  from scott.user_log
                  order by account ) loop

        if ( rec.stamp > sysdate - 90 ) then

            p_report := ... -- do something special

        else

            if ( rec.account = user ) then

                p_report := ... -- if current user then ...

            end if;

        end if;

        if ( rec.message like '%ABC%' ) then

            p_report := ... -- write a particular format

        else

            p_report := ... -- write another format

        end if;

    end loop;

end;
```

```

        end if;

    end loop;

end user_log_report;

--
view user_log_monthly
as
select stamp,
       account,
       message
  from scott.user_log
 where stamp between sysdate - 30
            and sysdate;

```

Assume we have already moved our data over to Hadoop and created a Hive table, of the same name. Our remote table looks like this ...

```

$ beeline -u jdbc:hive2://hive.corp.com:10000 \
-n oracle \
-w welcome1.passwd

0: jdbc:hive2://localhost:10000> desc user_log;
+-----+-----+-----+
| col_name | data_type | comment |
+-----+-----+-----+
| stamp    | date      |         |
| account  | string     |         |
| message  | string     |         |
+-----+-----+-----+
3 rows selected (0.17 seconds)

```

## Hive-ODCI configuration

Now that we have our table in Hadoop/Hive and we can connect via beeline, and see it there, let's setup Hive-ODCI as an access point to that table.

Because all of our clients will be accessing the same Hadoop/Hive data-store, we can setup a common connection strategy for Hive-ODCI via the parameters

```

dbms_hive.param( 'hive_jdbc_url', 'jdbc:hive2://hive.corp.com:10000' );
dbms_hive.param( 'hive_jdbc_url.1', 'user=oracle' );
dbms_hive.param( 'hive_jdbc_url.2', 'password=welcome1' );

```

If we have clients which need a different connection strategy or use different parameters, they can change these at the session level to meet their needs specifically, but for now we'll assume everyone is using the same thing.

## Hive-ODCI object creation

As denoted above we have PL/SQL code and a VIEW which is now invalid because the `SCOTT.USER_LOG` table no longer exists.

So, let's put it back using Hive-ODCI. First let's create a new VIEW which replaces the table.

```
--
grant execute on hive_q to scott;

--
create or replace view scott.user_log
(
    stamp,
    account,
    message
)
as
select *
from table( hive_q( q'[ select stamp,
                           account,
                           message
                           from user_log
                           order by stamp ]' ) )

/
```

Whew, that was easy. We now have a column-by-column replacement of the old table with a newly created remote table, that provides the same data types as before.

```
SQL> desc SCOTT.USER_LOG
```

| Name    | Null?    | Type           |
|---------|----------|----------------|
| STAMP   | NOT NULL | DATE           |
| ACCOUNT |          | VARCHAR2(4000) |
| MESSAGE |          | VARCHAR2(4000) |

Let's see if it worked ...

```
SQL> alter procedure scott.user_log_report compile;

Procedure altered.
```

Excellent, so let's take a look at replacing the view. In this case, we only care about data from the last 30 days, no need to make Hadoop/Hive do more work than it has to. Let's use bindings to restrict the data at the Hadoop/Hive layer instead of at the Oracle layer.

```
--
grant execute on hive_bind to scott;
grant execute on hive_binds to scott;
```

```
--
create or replace view scott.user_log_monthly
(
    stamp,
    account,
    message
)
as
select *
from table( hive_q( q'[ select stamp,
                        account,
                        message
                        from user_log
                        where stamp between ? and ? ]',
hive_binds( hive_bind( to_char( sysdate - 30,
                                'yyyy-mm-dd' ),
                        1 /* type_date */,
                        1 /* ref_in */ ),
hive_bind( to_char( sysdate, ,
                    'yyyy-mm-dd' ),
            1 /* type_date */,
            1 /* ref_in */ ) ) )
/
```

Let's break this one down, as it is a little more complex than the other. The `hive_q` function takes 3 parameters, 2 of which are defaulted to `NULL`. The first is a `hive_binds` object which is simply an array of `hive_bind` objects or individual bind data. The second parameter is the `hive_session` (e.g. URL, User, ...), but since we are using a common connection strategy we will ignore this argument, not pass it in, let it continue to be `NULL`.

The `hive_bind` object also takes 3 arguments, the first is the data to be used as the bind. The second is the type of data to be bound (e.g. string, date, number, etc...). And the third is the scope of the bind, most useful for DML operations (e.g. IN, OUT, IN/OUT).

In our case we only care about scope references of IN, hence the `1 /* ref_in */`. Since both bind operators are `DATE` variables they both have `1 /* type_date */`.

The actual bound data is based on `SYSDATE`, but since we don't know up front the `NLS_DATE_FORMAT` setting for the session, we simply guarantee the format by wrapping it in a `TO_CHAR()` function and specify the format Hadoop/Hive is expecting.

The Oracle RBAC controls dictate who can SELECT from the VIEWS, just as before. So, if you have custom roles which need access to the VIEWS, they can be granted access in the same way.

```
grant select on scott.user_log to my_app_role;
grant select on scott.user_log_monthly to public;
```

Now wala! We have a VIEW that will retrieve the last 30 days worth of data from Hadoop/Hive.

## Almost there

We now have replacements for both the PL/SQL Reporting and the monthly VIEW , but we're just not yet at a 100%, so what have we forgot?

You guessed it; DML. Our application is still putting data in, modifying it and removing it so we need to support that too. No problem. We can use an `INSTEAD OF` trigger and send our `INSERT`, `UPDATE` and `DELETE` commands to Hadoop/Hive using Hive-ODCI.

```

create or replace trigger scott.user_log_dml

instead of delete or insert or update on scott.user_log
for each row

declare

cmd varchar2( 4000 );
bnd hive_binds := hive_binds();

begin

if ( inserting ) then

cmd := q'[ insert into user_log
              ( stamp, account, message )
            values
              ( ?, ?, ? ) ]';

bnd.extend;
bnd( bnd.count ) := hive_bind( to_char( :new.stamp,
                                         'yyyy-mm-dd' ),
                               hive_binding.type_date,
                               hive_binding.ref_in );

bnd.extend;
bnd( bnd.count ) := hive_bind( :new.account,
                               hive_binding.type_string,
                               hive_binding.ref_in );

bnd.extend;
bnd( bnd.count ) := hive_bind( :new.message,
                               hive_binding.type_string,
                               hive_binding.ref_in );

elsif ( updating ) then

cmd := q'[ update user_log
              set account = ?,
                message = ?
            where stamp = ? ]';

bnd.extend;
bnd( bnd.count ) := hive_bind( :new.account,
                               hive_binding.type_string,
                               hive_binding.ref_in );

```



```

        bnd.extend;
        bnd( bnd.count ) := hive_bind( :new.message,
                                         hive_binding.type_string,
                                         hive_binding.ref_in );

        bnd.extend;
        bnd( bnd.count ) := hive_bind( to_char( :new.stamp,
                                                  'yyyy-mm-dd' ),
                                         hive_binding.type_date,
                                         hive_binding.ref_in );

    elsif ( deleting ) then

        cmd := q'[ delete from user_log
                    where stamp = ? ]';

        bnd.extend;
        bnd( bnd.count ) := hive_bind( to_char( :new.stamp,
                                                  'yyyy-mm-dd' ),
                                         hive_binding.type_date,
                                         hive_binding.ref_in );

    else

        null; -- should never get here

    end if;

    if ( ( cmd is not null ) and ( bnd.count > 0 ) ) then

        -- execute the remote statement
        hive_remote.dml( cmd, bnd );

    end if;

end user_log_dml;
/

```

This may look more complex, but it's really not. We are simply creating a trigger to handle the DML events and passing them off to Hive-ODCI. The `HIVE_REMOTE.DML()` works in the same way as a query, accepting an array of bind objects and connection information.

The only real difference here is the example is creating and using local variables for `HIVE_BINDS` and the `HIVE_BINDING` types to show how it can be used in that manner.

## Hints

The Hive SQL-Like for Hadoop allows for hints to be passed in which control plan execution and join optimization, the Hive-ODCI interface also allows for hints to be passed to the ODCI engine which turn on or provide additional functionality. These hints are parsed and removed from the query before being passed off to Hive for processing. Any hints not supported by Hive-ODCI remain in place and are passed in as provided.

All hints are provided as comment blocks to the query, just like an Oracle or Hive hint. Comment blocks are wrapped, starting with `/*+` and ending in `*/` or start with `--+` and end with a newline.

### **typecast()**

This hint provides a mechanism for casting data types in Hive to ones provided to Oracle. This is useful when the Hive data type is not supported by Oracle or the data type in Hive is insufficient for describing the Oracle type.

Take for example the universal Hive data type `string`, which by default is handled as a `VARCHAR2(4000)` by Hive-ODCI. However, suppose that the original data type was a `CLOB`, or simply that the `string` column exceeds the `4000` character limit? This is where the `typecast()` hint comes into play. You can instruct Hive-ODCI to create a `CLOB` column data type when the column name is encountered with a casting rule.

Casting rules are either space delineated while wrapped in the parenthesis of the `typecast()` keyword. Each rule is formatted by the column name in the Hive record-set delineated by a colon with the following data type and optionally length, precision and scale.

```
column_name:datatype[(length/precision,scale)]
```

Let's take the above SQL and assume that when Scoop was used to originally copied to HDFS that it was done by an over zealous administrator, who took the high-road and made every Hive column a `string`. Oops, that's not going to work for us because the original table had the `STAMP` column as a `DATE` and the `MESSAGE` column as a `CLOB` and we know that it will exceed `4000` characters in length.

So let's confirm by first describing the table ...

```
SQL> col "col_name" for a30
SQL> col "data_type" for a30
SQL> col "comment" for a20

SQL> select * from table( hive_q( 'desc user_log', null, null ) );

col_name      data_type      comment
-----
stamp         string
account       string
message       string
```

If we cast those types to the correct data types, using `typecast()` we can rest assure that the Oracle view or cursor returned is the correct data type.

If we recreate or view, this time using the casting rule, as follows

the type defined in Oracle will now be correct.

```
--
create or replace view scott.user_log
(
    stamp,
    account,
    message
)
as
select *
from table( hive_q( q'[ select /*+ stamp:date message:clob */
                        stamp,
                        account,
                        message
                        from user_log
                        order by stamp ]' ) )
/
```

Now when we look at our **VIEW** it is described correctly ...

```
SQL> desc SCOTT.USER_LOG
```

| Name    | Null? | Type         |
|---------|-------|--------------|
| STAMP   |       | DATE         |
| ACCOUNT |       | VARCHAR2(30) |
| MESSAGE |       | CLOB         |

For other types, you can optionally define length, precision and scale.  
Some other examples may look like ...

```
select /*+
    cust_id:number(9)
    salary:number(7,2)
    dob:timestamp(2)
    first_name:varchar2(50)
    last_name:varchar2(100)
    middle_initial:char(1)
*/
from customers
```

If you have a large number of columns, you can use the helper  
function **HIVE\_HINT()** (below), against a template, to generate the hints  
for you.

```
SQL> col hint for a80 word_wrap
SQL> select hive_hint( 'SCOTT', 'CUST' ) hint from dual;

HINT
-----
/*+ cust_id:number(9) salary:number(7,2)
dob:timestamp(2) first_name:varchar2(50)
last_name:varchar2(100)
middle_initial:char(1) */
```

If the column is not encountered during execution, Hive-ODCI simply ignores the rule.

## Final thought

Note we have **not** changed any code in our application or in our PL/SQL procedure. Everything remains exactly as it was before, but our data exists only in Hadoop/Hive. We can view the data, use it with our PL/SQL and even manipulate it with DML. So, you can go to your meeting now and be the hero.

# Guidelines

---

The following sections are guidelines based on practical real-world experience and are intended to help you in the decision making process only. They are not hard-and-fast rules of "thou shalt not" commandments that must be followed to use Hive-ODCI.

The sections are separated into focus areas for Developers and Administrators. In the real-world these are often blurry, completely undefined and are overlapping. While that is true, no matter what role you play, the distinction between them should always be upheld as much as possible.

## For Developers

---

While Hive-ODCI is intended to help the transition of moving data into a Hadoop/Hive data-store, it can with planning and understanding make it transparent to your customers. But at some point it may require Developer intervention.

## Performance Considerations

Hive-ODCI is a `PIPELINE` type, so its performance is impacted almost entirely on the performance of Hadoop/Hive. So make sure that Hive is performing at optimal levels.

Ensure that when moving data from Oracle to Hadoop/Hive or creating new tables you have designed in performance from the beginning.

## Analytics over in-line views

If you are using in-line views within the query, you may find that using the built-in `RANK()` and `OVER()` functions perform much better.

For example, take the following query

```
select user_log.*
  from user_log join
    ( select account,
      max( stamp ) as stamp
```

```
        from user_log
        group by account ) inline
on user_log.account = inline.account
and user_log.stamp = inline.stamp;
```

While there is nothing technically wrong with this query, based on the plan it could be much better. let's try again ...

```
select *
from ( select *,
        rank() over
          ( partition by account,
            order by stamp ) as rank
        from user_log ) ranking
where ranking.rank = 1;
```

Much better. We get the same result with almost a magnitude in increase on performance. If all of this looks familiar, then your right. This has been in Oracle RDBMS for years, so use the same common-sense techniques in your Hadoop/Hive queries as you would in your Oracle RDBMS queries.

## The CBO

In recent additions of Hadoop/Hive a Cost-Based-Optimizer (CBO) was introduced, which like in Oracle RDBMS, performs optimizations based on cost, which can adjust the execution plan based on order joins, types of joins, degree of parallelism, etc... leading to increases in query performance.

If not enabled globally, you can still use the Hadoop/Hive CBO, by setting the following parameters at the beginning of the query.

```
set hive.cbo.enable=true;
set hive.compute.query.using.stats=true;
set hive.stats.fetch.column.stats=true;
set hive.stats.fetch.partition.stats=true;
```

Just like Oracle RDBMS, it's important to analyze the tables so the CBO has current cost information. For example, collect statistics at the table and column levels as necessary.

```
analyze table my_table compute statistics;
analyze table my_table compute statistics for columns;
analyze table my_table compute statistics for columns col1, col2;
```

## ORCFile

Using ORCFile for Hadoop/Hive table should already be a matter of practice because it is shown to be extremely beneficial in getting fast response times for queries.

If existing tables are not already ORC then it would be prudent to migrate them as soon as possible, even if you convert them by-hand.

However, if possible the best case would be to modify the ingest process to use ORC up front.

## Apache Tez

Whenever possible use the Apache Tez execution engine instead of a Map-reduce engine. If it is not enabled by default in your environment, then use the following setting at the beginning of the query.

```
set hive.execution.engine=tez;
```

## Vectorized queries

Like scans, aggregations, filters and joins vectorized query execution improves performance of operations by splitting them into batches of 1024 rows at a time instead of single row. If not enabled for your environment then use the following setting to at the beginning of the query.

```
set hive.vectorized.execution.enabled = true;  
set hive.vectorized.execution.reduce.enabled = true;
```

## New horizons

Putting Hive-ODCI into practice is not all that hard, but it is a new way of doing business. And with "new" things come hiccups and unforeseen circumstances that can put you behind or stop you dead in your tracks.

The biggest thing here, is to identify those risks upfront and mitigate them as soon as possible. Don't wait until the last minute, plan ahead and you won't regret it. I know this is "touchy-feely" advice that you can get at the local coffee shop, but we often forget about the small things in our rush.

## Change your PL/SQL code

If your PL/SQL can be changed for the better when using Hive-ODCI, do it. In reviewing your PL/SQL you may find that leveraging the functionality provided by Hive-ODCI makes it more readable, faster, maintainable, etc... then go ahead and make arrangements to modify the existing code base or introduce new code.

Sometimes, that's not feasible or even possible but many times we (the big we) have a tendency to force-fit a solution when that isn't warranted based on the situation.

## Keep signatures consistent

Even with the above statement being true, it's best to keep the underlying `TABLE` and `VIEW` objects with the same signature, column names, types, lengths, etc...

## Familiarize yourself with the API

Finally, make sure you familiarize yourself and are comfortable with the Hive-ODCI API. I know reading documentation is boring (writing it even more so, trust me) but take the time and go through it as many times as needed to get a firm understanding. This is particularly important for Developers, as they are the "real" users of the system.

Write some test scripts, see how it behaves in particular situations and make darn sure you know what's coming when you start inserting this stuff into your process.

If you've read the manual, tried it out, found a bug, or simply do not understand, then see the Authors section for POC information and contact me. I'll be happy to help where I can or provide advice and moral support as needed.

## For Administrators

---

Administrators are by enlarge the last line of defense against all manner of issues, they are the gatekeepers for security, performance, storage, new technologies, legacy systems, and a full onslaught of change management from every direction.

Whether the systems are large or small, clustered or single instance the administrator's toolbox is typically a plethora of knowledge, scripts and documentation. Hive-ODCI attempts add to that toolbox, without being just "something else to learn" by leveraging common facilities which already exist for the administrator. Most Hive-ODCI footprint objects will be first-class citizens in the Oracle ecosystem and will be managed, secured and monitored in the same way as other object types.

## Wait events

An administrator knows which wait events to look for and which ones can be ignored or are simply part of a running system. Hive-ODCI does not change those presumptions. Mainly because the load Hive-ODCI incurs is on the remote Hadoop/Hive system and not the Oracle RDBMS.

Something to be aware of, however is that unless all your users are hitting Hive-ODCI object simultaneously a wait event for Hive-ODCI will not float to the top of the wait percentages.

A Hive-ODCI object having issue would still be in one of the more innocuous wait events or classes. Additionally, it won't be in a blocking state, so you have to recognize those events which may be indicating problems for the clients.

## WAIT\_TIME and SECONDS\_IN\_WAIT

In the `GV$SESSION` view the columns `WAIT_TIME` and `SECONDS_IN_WAIT` can easily be used to determine if Hive-ODCI is waiting on a remote

call to respond or complete.

## SQL\*Net message to client

This wait event can be observed while the Hive-ODCI is sending or responding to the call. While this event itself is not an issue, one that is taking a longer time than normal may indicate that the `query_limit` parameter is set too high.

## buffer latch and latch free

While these in normal circumstances, indicate query issues they do not necessarily mean the same thing for Hive-ODCI. You may get these events when the Hive-ODCI is materializing large amounts of data, for example in local sorting or aggregate processing of the returned data.

If this event is observed too high, or too often for a particular Hive-ODCI object then look into off-loading the sort, group or aggregate operations to Hadoop/Hive instead of locally.

Take for example the following

```
select *
  from table( hive_q( q'[ select cust_id,
                             last_name,
                             first_name
                             from cust ]',
                             null,
                             null ) )
 order by cust_id;
```

If this query is showing up in a wait event that is indicating issues materializing the `cust_id` column locally for sorting then considered making the following change to off-load that operation to Hadoop/Hive

```
select *
  from table( hive_q( q'[ select cust_id,
                             last_name,
                             first_name
                             from cust
                             order by cust_id ]',
                             null,
                             null ) );
```

This produces the same end result, but without local materialization and with less wait activity in the Oracle RDBMS.

## Storage

Hive-ODCI is a zero storage object solution, like a `VIEW`, for



the majority of how its used. However certain aspects of Hive-ODCI will consume space, such as the Hive-ODCI Log and the Saved Filters (Binding).

It is unlikely that your user community, no matter how large, or how much Hive-ODCI is used will generate Filters that impact your storage, necessitating TS extents beyond what was allocated during the installation.

The Hive-ODCI Log data is another story. Depending on the `log_level` set for the System, large amounts of data can be generated. The `log_level` value is a bitmask, which turns types on and off. The types are detailed below but also here

```
create or replace package impl as

--
  none constant number := 0;
  error constant number := 1;
  warn constant number := 2;
  info constant number := 4;
  trace constant number := 8;

  ...

end impl;
```

Each level is progressively verbose with what is written. The types should be obvious, but needless to say `error` writes only critical exceptions, while `trace` writes all operations. So a value of `3` would be `error + warn` and a value of `31` would be `error + warn + info + trace`.

If you find that the Hive Log is filling up faster than expected, review the column `VALUE` in the `DBA_HIVE_PARAMS` to determine the current value.

```
select value
  from dba_hive_params
 where name = 'log_level';
```

If you find that the `log_level` is set appropriately, or as you expected, then this means that the client has set the `session_log_level()` to something too high, which may need to be changed or justified.

The account which created the log data is stored in the column `NAME` found in the `DBA_HIVE_LOG` view. Use the following to get counts of the name generating the most logging information.

```
select name,
       count(0) total
```

```
from dba_hive_log
group by name
order by 2 desc;
```

Log information can be purged by any trusted account granted the `HIVE_ADMIN` role using the following statement.

```
exec dbms_hive.purge_log;
```

Be aware that the statement removes **ALL** log data.

## Move to a different TS

By default Hive-ODCI installs to the same tablespace assigned to the `SYSTEM` schema. This for most installations will be sufficient. If you disagree with the installation choice, or you have a larger storage consumption than anticipated you can move Hive-ODCI to another tablespace using the `MOVE_TS` procedure in the `DBMS_HIVE` package.

```
create or replace package dbms_hive as

...

--
procedure move_ts( ts in varchar2,
                  obj in varchar2 default null );

end dbms_hive;
```

## Purge log and filters

If the Hive-ODCI log data or filter data becomes too large, it can easily be purged using the `PURGE_LOG` and `PURGE_FILTER` procedures respectively in the `DBMS_HIVE` package.

```
create or replace package dbms_hive as

...

--
procedure purge_log;
procedure purge_filter( key in varchar2 default null );

...

end dbms_hive;
```

These procedures do exactly what they indicate. And because they are exposed through the Hive-ODCI management package they can be called by a trusted user who has been granted the `HIVE_ADMIN` role.

# Role Based Access Control

As indicated multiple times now, Hive-ODCI objects are primarily first-class and are managed with RBAC just like any other object.

When Hive-ODCI is installed, it creates 2 helper roles for you setting up the appropriate permissions for their use.

The `HIVE_USER` role has access to the common Hive-ODCI objects making it possible to Query, use DML or DDL, manage Bind variables, etc...

The second role, is `HIVE_ADMIN` which provides access to the management of Hive-ODCI itself, such as Logs, Parameters, etc... This role is also granted to the `DBA` role by default so be cognizant that a user with `HIVE_ADMIN` is considered a trusted and responsible party, similar to a user with `DBA` but having less permissions in the database.

## Privileges

Hive-ODCI has built-in RBAC controls for securing the rights of remote object access. This works similar to the Oracle RBAC grant/revoke capabilities on database objects.

The Hive-ODCI privileges allow management of user and role access assignments through the granting and revoking of permissions on remote table names. This is controlled by the API available in the `DBMS_HIVE` package procedures `GRANT_ACCESS` and `REVOKE_ACCESS`.

SQL command operations submitted for execution in Hive/HDFS are preprocessed by the built-in security control interface, which extracts the table name list and checks it against the authorization assignment. Because Hive-ODCI cannot predetermine schema and table names without executing the command, which would be in contradiction to providing the security in the first place, the API for granting and revoking privileges are available only in the `DBMS_HIVE` package, requiring `HIVE_ADMIN` access.

```
create or replace package dbms_hive as

...

procedure grant_access( opr in varchar2,    -- operation list
                      tab in varchar2,    -- hive table list (case sensitive)
                      gnt in varchar2 ); -- grantee list

-- revoke access from hive table
procedure revoke_access( opr in varchar2,
                       tab in varchar2,
                       gnt in varchar2 );

...
```

```
end dbms_hive;
```

Each parameter, `opr` (operation), `gnt` (grantee) and `tab` (table) can be provided as a comma separated list of items, rather than one at a time thus making the calls more convenient in their use.

Role grants are always verified based hierarchy, so that grants of role-to-role can still provide access. For example, if a role named `MANAGER` is granted to a role named `WORKER`, and the `WORKER` is granted Hive-ODCI `SELECT` privileges on a table, then anyone with the `MANAGER` role can still gain query access, even if the role `WORKER` was not directly granted to the user.

Examples of usage:

```
-- grant single operation, on a single table to a single user
SQL> exec dbms_hive.grant_access( 'select', 'cust', 'scott' );

-- grant multiple operations, on a single table to a single role
SQL> exec dbms_hive.grant_access( 'select, insert', 'cust', 'dba' );

-- grant multiple operations, on multiple tables to multiple users and role;
SQL> exec dbms_hive.grant_access( 'select, insert, update, delete', -
2                                'cust, movie, review', -
3                                'scott, jdoe, dba' );

-- revoke an operation, on a table from multiple users and roles
SQL> exec dbms_hive.revoke_access( 'delete', 'review', 'scott, jdoe' );
```

Access assignments can be seen through the provided views `DBA_HIVE_PRIVS` and `USER_HIVE_PRIVS`. The latter is available only with `HIVE_ADMIN` access and the former through `HIVE_USER` access. These views will report the privilege granted for the operation, even if it exists lower in the hierarchical layer. This is useful for determine how or if a user is able to submit an operation.

As the run-time functionality, parses and makes privilege determinations, the Hive-ODCI security uses multiple parameters to determine action and outcome of resulting calls. The parameters `auth_no_grant_action`, `auth_auto_grant`, `auth_table_undefined` and `auth_sql_parse_error` drive what Hive-ODCI does as a result of the security validations and assignments.

## Bindings

Filters, or saved `HIVE_BINDS` work more like Network ACL (see `DBMS_NETWORK_ACL_ADMIN`) and Java Policies (see `DBMS_JAVA`) than they do first-class object permissions.

When created a Filter will by default assign ownership of the `HIVE_BINDS` object key to the account which created it.

That account, or a `HIVE_ADMIN` assigned account can grant and revoke permissions to both users and roles using the

`HIVE_BINDING` type methods.

```
create or replace package binding as

...

--
priv_read      constant guard    := 1;
priv_write     constant guard    := 2;
priv_readwrite constant guard    := 3;

--
procedure allow( key in varchar2,
                 act in varchar2,
                 lvl in guard default priv_readwrite );

--
procedure deny( key in varchar2,
                act in varchar2 );

...

end binding;
```

The `ALLOW()` procedure works like a `GRANT` and the `DENY()` works like a `REVOKE`

When a `HIVE_BIND` is accessed from the stored Filter the permission is checked for the operation (e.g. read, write, etc...). If a `PUBLIC` grant is made, then all accounts may access the Filter at the level granted.

## Java Policies

Hive-ODCI will in no way automatically grant or revoke Java Policy permissions, therefore it is necessary for the administrator to take on this action as needed.

When Hive-ODCI attempts to set a Java Property at run-time, which has not been granted the appropriate policy an exception is thrown. As part of that exception, the necessary `DBMS_JAVA.GRANT_PERMISSION()` is provided to the caller as part of the exception text. It is likely that the text generated by the database will be sent to the administrator requesting it be executed.

Before doing so, make sure there is a clear understanding of what that permission provides as well as what the property enables or achieves.

Failure to understand the property and/or permission may result in reducing the security posture of the database or at worst make the database unstable. For the most part this will not happen, but it is good practice to have full situational awareness before proceeding.

## API

---

The Application programming Interface (API) for Hive-ODCI is accessible through the PL/SQL objects created during installation. The objects include Views, Packages, Procedure, Types and Functions each providing a unique set of functionality based on need.

## Views

---

Views are the insight into Hive-ODCI, providing invaluable information for both Administrators and Users.

### dba\_hive\_params

Lists both the system and session level values.

This view is accessible only by the `HIVE_ADMIN` role

- Columns

|               |                                |
|---------------|--------------------------------|
| name          | - Parameter name               |
| session_value | - Current session <b>value</b> |
| system_value  | - Default system <b>value</b>  |

### dba\_hive\_filters

Lists the saved Filters.

This view is accessible only by the `HIVE_ADMIN` role

- Columns

|             |                                     |
|-------------|-------------------------------------|
| key         | - Filter key name                   |
| seq         | - The ordinal <b>bind</b> sequence  |
| <b>type</b> | - The <b>bind type</b>              |
| scope       | - The <b>bind</b> reference scope   |
| value       | - The value of the <b>bind</b> data |

### dba\_hive\_filter\_privs

Lists the privileges of the saved Filters.

This view is accessible only by the `HIVE_ADMIN` role

- Columns

|              |   |
|--------------|---|
| key          | - Filter key name                       |
| grantee      | - The grantee name, user <b>or</b> role |
| <b>read</b>  | - Contains <b>read</b> access           |
| <b>write</b> | - Contains <b>write</b> access          |

### dba\_hive\_privs

Lists all privilege assignments for the remote Hive/HDFS operation rights.

This view is accessible only by the `HIVE_ADMIN` role

- Columns

|                           |  |
|---------------------------|--|
| <code>table_name</code>   | - The Hive/HDFS table name (case sensitive) being protected                    |
| <code>privilege</code>    | - The protected operation  |
| <code>grantee</code>      | - The user <b>or</b> role assigned rights to the operation <b>of</b> the table |
| <code>grantee_type</code> | - The type (USER <b>or</b> ROLE) <b>of</b> the grantee                         |

## dba\_hive\_log

Lists the log data

This view is accessible only by the `HIVE_ADMIN` role

- Columns

|                    |  |
|--------------------|--|
| <code>stamp</code> | - The time stamp the log record                        |
| <code>name</code>  | - Account name who <b>create</b> the <b>log</b> record |
| <code>type</code>  | - The type <b>of log</b> record                        |
| <code>tier</code>  | - The <b>text</b> representation <b>of</b> the type    |
| <code>text</code>  | - The <b>text of the log</b> record                    |

## user\_hive\_log

Lists the log data for the current user

This view is accessible through the `HIVE_USER` role

- Columns

|                    |   |
|--------------------|---|
| <code>stamp</code> | - The <b>time</b> stamp the <b>log</b> record |
| <code>type</code>  | - The type of <b>log</b> record               |
| <code>tier</code>  | - The text representation of the type         |
| <code>text</code>  | - The text of the <b>log</b> record           |

## user\_hive\_params

Lists current parameter values used by the session

This view is accessible by both the `HIVE_ADMIN` and `HIVE_USER` roles

- Columns

|                    |                          |
|--------------------|--------------------------|
| <code>name</code>  | - Parameter name         |
| <code>value</code> | - Parameter <b>value</b> |

## user\_hive\_filters

Lists the saved Filters which are accessible by the session

This view is accessible by both the `HIVE_ADMIN` and `HIVE_USER` roles

- Columns

|       |   |   |
|-------|---|---|
| key   | - | Filter key name                         |
| seq   | - | The ordinal <code>bind</code> sequence  |
| type  | - | The <code>bind type</code>              |
| scope | - | The <code>bind</code> reference scope   |
| value | - | The value of the <code>bind</code> data |

## user\_hive\_filter\_privs

Lists the privileges of the saved Filters which are accessible by the session

This view is accessible by both the `HIVE_ADMIN` and `HIVE_USER` roles

- Columns

|              |   |                                       |
|--------------|---|---------------------------------------|
| key          | - | Filter key name                       |
| grantee      | - | The grantee name, user <b>or</b> role |
| <b>read</b>  | - | Contains <b>read</b> access           |
| <b>write</b> | - | Contains <b>write</b> access          |

## user\_hive\_privs

Lists the currently logged in user's privilege assignments for the remote Hive/HDFS operation rights.

This view is accessible by both the `HIVE_ADMIN` and `HIVE_USER` roles

- Columns

|              |   |  |
|--------------|---|--|
| table_name   | - | The Hive/HDFS table name (case sensitive) being protected                      |
| privilege    | - | The protected operation  |
| grantee      | - | The role assigned rights providing access to the operation <b>of</b> the table |
| grantee_type | - | The <b>ROLE</b> <b>of</b> the grantee  |

## Packages

---

### binding

An interface for creating, manipulating and saving `HIVE_BINDS`



arrays of `HIVE_BIND` objects. These objects are used in the marshaling and interpretation of arguments to and from the JDBC Driver providing both data type and direction.

This object is accessible through the `HIVE_USER` role and is referenced as the `SYNONYM` name `HIVE_BINDING`

## Subtypes

- reference -- `number` for bind scope values
- typeof -- `number` for bind value types
- guard -- `number` for bind access privileges

## Constants

Helper constants for abstracting type information

### Generic

| Name              | Value          |
|-------------------|----------------|
| -----             |                |
| <code>none</code> | <code>0</code> |

### Privilege

| Name                        | Value          |
|-----------------------------|----------------|
| -----                       |                |
| <code>priv_read</code>      | <code>1</code> |
| <code>priv_write</code>     | <code>2</code> |
| <code>priv_readwrite</code> | <code>3</code> |

### Scope

| Name                     | Value          |
|--------------------------|----------------|
| -----                    |                |
| <code>scope_in</code>    | <code>1</code> |
| <code>scope_out</code>   | <code>2</code> |
| <code>scope_inout</code> | <code>3</code> |

### Type

| Name                        | Value           |
|-----------------------------|-----------------|
| -----                       |                 |
| <code>type_bool</code>      | <code>1</code>  |
| <code>type_date</code>      | <code>2</code>  |
| <code>type_float</code>     | <code>3</code>  |
| <code>type_int</code>       | <code>4</code>  |
| <code>type_long</code>      | <code>5</code>  |
| <code>type_null</code>      | <code>6</code>  |
| <code>type_rowid</code>     | <code>7</code>  |
| <code>type_short</code>     | <code>8</code>  |
| <code>type_string</code>    | <code>9</code>  |
| <code>type_time</code>      | <code>10</code> |
| <code>type_timestamp</code> | <code>11</code> |
| <code>type_url</code>       | <code>12</code> |

## Interface

### get()

Overloaded **FUNCTION** to retrieve a saved **HIVE\_BINDS** filter based on Key name or returns a single **HIVE\_BIND** from the session array.

- Prototype

```
function get( key in varchar2 ) return binds;  
  
function get( idx in number, lst in binds ) return bind;
```

- Parameter

|     |   |                             |
|-----|---|-----------------------------|
| key | - | Saved key name              |
| idx | - | Ordinal number of the item. |

### count()

Overloaded **FUNCTION** to count the length of a saved **HIVE\_BINDS** filter based on Key name or an existing one provided.

- Prototype

```
function count( key in varchar2 ) return number;  
function count( lst in binds ) return number;
```

- Parameter

|     |   |                           |
|-----|---|---------------------------|
| key | - | Saved key name            |
| lst | - | Existing HIVE_BINDS array |

### new()

A **FUNCTION** to create a new, single, **HIVE\_BIND** object.

- Prototype

```
function new( value in varchar2,  
             type in typeof default type_string,  
             scope in reference default scope_in ) return bind;
```

- Parameter

|       |   |   |
|-------|---|---|
| value | - | Value of the HIVE_BIND                              |
| type  | - | Interpretation type of the HIVE_BIND                |
| scope | - | Reference scope, direction, for using the HIVE_BIND |

### append()

Overloaded `PROCEDURE` to append a `HIVE_BIND` object to an existing `HIVE_BINDS` array, saved or passed in.

- Prototype

```
procedure append( key    in varchar2,  
                  value in varchar2,  
                  type  in typeof    default type_string,  
                  scope in reference default scope_in );  
  
procedure append( value in    varchar2,  
                  type  in    typeof    default type_string,  
                  scope in    reference default scope_in,  
                  lst   in out binds );  
  
procedure append( key in varchar2, val in bind );  
  
procedure append( val in bind, lst in out binds );  
  
procedure append( key in varchar2, val in binds );  
  
procedure append( val in binds, lst in out binds );
```

- Parameter

|       |   |  |
|-------|---|--|
| key   | - | Saved key name   |
| value | - | Value of the <code>HIVE_BIND</code>  |
| type  | - | Interpretation type of the <code>HIVE_BIND</code>  |
| scope | - | Reference scope, direction, <b>for using</b> the <code>HIVE_BIND</code>                              |
| lst   | - | Existing <code>HIVE_BINDS</code> <a href="#">array</a>   |
| val   | - | A <code>HIVE_BINDS</code> <a href="#">array</a> to be appended to the existing <a href="#">array</a> |

### **change()**

A `PROCEDURE` to modify an exiting element of a saved `HIVE_BIND` object.

- Prototype

```
procedure change( key    in varchar2,  
                  idx   in number,  
                  value in varchar2,  
                  type  in typeof    default type_string,  
                  scope in reference default scope_in );
```

- Parameter

|       |   |   |
|-------|---|---|
| key   | - | Saved key name <b>of</b> the array                                      |
| idx   | - | Ordinal number <b>of</b> the item to modify                             |
| value | - | Value <b>of</b> the <code>HIVE_BIND</code>                              |
| type  | - | Interpretation type <b>of</b> the <code>HIVE_BIND</code>                |
| scope | - | Reference scope, direction, <b>for</b> using the <code>HIVE_BIND</code> |

### **remove()**

A **PROCEDURE** to delete an exiting element of a saved **HIVE\_BIND** object.

- Prototype

```
procedure remove( key in varchar2,  
                  idx in number );
```

- Parameter

|     |   |                                      |
|-----|---|--------------------------------------|
| key | - | Saved key name of the array          |
| idx | - | Ordinal number of the item to modify |

#### **replace()**

A **PROCEDURE** to replace one **HIVE\_BINDS** array with another.

- Prototype

```
procedure replace( val in binds, lst in out binds );
```

- Parameter

|     |   |   |
|-----|---|---|
| val | - | A <b>HIVE_BINDS array</b> to be appended to the existing <b>array</b> |
| lst | - | <b>HIVE_BINDS array</b> to overwrite                                  |

#### **clear()**

An overloaded **PROCEDURE** to remove all elements in a saved Filter or a passed in array.

- Prototype

```
procedure clear( key in varchar2 );  
procedure clear( lst in out binds );
```

- Parameter

|     |   |                                      |
|-----|---|--------------------------------------|
| key | - | Saved key name of the <b>array</b>   |
| lst | - | <b>HIVE_BINDS array</b> to overwrite |

#### **allow()**

A **PROCEDURE** to grant access rights to saved Filter. If an existing grant was already made, then it is replaced by the change

- Prototype

```
procedure allow( key in varchar2,  
                act in varchar2,  
                lvl in guard default priv_readwrite );
```

- Parameter

|     |   |  |
|-----|---|--|
| key | - | Saved key name <b>of</b> the array   |
| act | - | Name <b>of</b> the account to allow access, which can be an Oracle username <b>or</b> role |
| lvl | - | The level <b>of</b> access to grant  |

#### deny()

A **PROCEDURE** to revoke access rights from a saved Filter. If no existing grant exists an exception is thrown.

- Prototype

```
procedure deny( key in varchar2,
               act in varchar2 );
```

- Parameter

|     |   |  |
|-----|---|--|
| key | - | Saved key name <b>of</b> the array   |
| act | - | Name <b>of</b> the account to allow access, which can be an Oracle username <b>or</b> role |

#### save()

A **PROCEDURE** to save a passed in **HIVE\_BINDS** array. If a key of the same name already exists, an exception is thrown.

- Prototype

```
procedure save( key in varchar2, lst in binds );
```

- Parameter

|     |   |                                    |
|-----|---|------------------------------------|
| key | - | Saved key name of the <b>array</b> |
| lst | - | <b>HIVE_BINDS array</b> to save    |

## Exception

### ex\_unknown

This exception is thrown when an unknown error has been encountered.

```
ex_unknown exception;
es_unknown constant varchar2( 256 ) := 'Unknown error encountered';
ec_unknown constant number := -20001;
pragma exception_init( ex_unknown, -20001 );
```

### ex\_denied

This exception is thrown when a saved Filter access attempt has failed

because of insufficient privileges.

```
ex_denied    exception;
es_denied    constant varchar2( 256 ) := 'Request denied,
                                         insufficient privileges';

ec_denied    constant number := -20002;
pragma       exception_init( ex_denied, -20002 );
```

## ex\_no\_grant

This is exception is thrown when a revoke attempt is made on a saved Filter that does not contain a previous grant for the account specified

```
ex_no_grant  exception;
es_no_grant  constant varchar2( 256 ) := 'Privileges not granted';
ec_no_grant  constant number := -20003;
pragma       exception_init( ex_no_grant, -20003 );
```

## dbms\_hive

A `PACKAGE` providing the interfaces for management of Hive-ODCI configurations.

This object is accessible through the `HIVE_ADMIN` role and is referenced as the `SYNONYM` name `DBMS_HIVE`

### Interface

#### exist()

A `FUNCTION` returning a `BOOLEAN` value  
( `true` or `false` ) if the parameter name exists

- Prototype

```
function exist( name in varchar2 ) return boolean;
```

- Parameter

```
name      -   Case sensitive name of the parameter to check
```

#### param()

An overloaded `FUNCTION` and `PROCEDURE` to both set and get a parameter. When setting, if the parameter exists then it's value is overwritten

- Prototype

```
-- get
function param( name in varchar2 ) return varchar2;

-- set
procedure param( name in varchar2, value in varchar2 );
```

- Parameter

|       |   |  |
|-------|---|--|
| name  | - | Case sensitive name of the parameter to get or <b>set</b>  |
| value | - | <b>Value of</b> the parameter <b>when</b> setting. The <b>value</b> is the <b>return for</b> the <b>FUNCTION</b> |

If the a parameter does not exist, then the function returns `NULL`

#### **remove()**

A `PROCEDURE` for removing a parameter. If the parameter does not exist, no error is thrown.

- Prototype

```
procedure remove( name in varchar2 );
```

- Parameter

|      |   |  |
|------|---|--|
| name | - | Case sensitive name <b>of</b> the parameter to <b>delete</b> |
|------|---|--|

#### **purge\_log()**

A `PROCEDURE` which purges all data in the `DBA_HIVE_LOG` optionally removing single user records. The `name` parameter is case-sensitive and when provided removes records matching only the those records created by `name`. If no `name` is provided or `NULL` is provided the table is truncated and storage dropped ( `TRUNCATE TABLE LOG$ DROP STORAGE` ). This is to provide a faster cleanup procedure than issuing an unqualified `DELETE` statement.

- Prototype

```
procedure purge_log( name in varchar2 default null );
```

#### **purge\_filter()**

A `PROCEDURE` for removing all saved Filters by name. If no name value or a `NULL` is provided as the argument then **ALL** Filters are removed.

- Prototype

```
procedure purge_filter( key in varchar2 default null );
```

- Parameter

|     |   |  |
|-----|---|--|
| key | - | Key name <b>of</b> the array to remove, <b>this</b> defaults |
|-----|---|--|

to NULL which indicates ALL values are removed

#### **move\_ts()**

A **PROCEDURE** for moving the Hive-ODCI objects to different tablespace. If the current tablespace name is the same as the one specified then no errors are thrown. If the tablespace specified does not exist then an exception is thrown.

An **OBJECT** name is alternatively provided which will move only a single object. If no value or **NULL** is used then **ALL** objects are moved.

Only the values " **NULL** ", "param", "filter", "priv" or "log" can be used, all other values will be ignored.

- **Prototype**

```
procedure move_ts( ts in varchar2,  
                  obj in varchar2 default null );
```

- **Parameter**

|     |   |  |
|-----|---|--|
| ts  | - | Destination tablespace name.   |
| obj | - | Optional object name, can be <b>NULL</b> , "param", "filter", "priv" or "log". Any other value will be ignored |

#### **grant\_access**

**PROCEDURE** used for assigning privileged access to remote Hive/HDFS operations. This procedure allocates the availability of user rights on those objects. This is *not* RBAC at Hive/HDFS, but rather security assignment for privileges to submit operations.

- **Prototype**

```
procedure grant_access( opr in varchar2,  
                       tab in varchar2,  
                       gnt in varchar2 );
```

- **Parameter**

|     |   |  |
|-----|---|--|
| opr | - | operation list, such as <b>SELECT, INSERT, UPDATE, CREATE</b> , etc... |
| tab | - | <b>case</b> sensitive list of <b>table/view names</b> to protect       |
| gnt | - | list of grantee assignments  |

Each parameter is a comma separated list of values, so that multiple values can be provided in a single call.

#### **revoke\_access**

**PROCEDURE** used for removing privileged access on remote Hive/HDFS operations. This removes previously granted allocation assignment, using



the `grant_access()` API above.

- Prototype

```
procedure revoke_access( opr in varchar2,  
                        tab in varchar2,  
                        gnt in varchar2 );
```

- Parameter

|     |   |  |
|-----|---|--|
| opr | - | operation list, such as <b>SELECT, INSERT, UPDATE, CREATE</b> , etc... |
| tab | - | <b>case</b> sensitive list of <b>table/view</b> names                  |
| gnt | - | list of <b>current</b> grantee assignments                             |

Each parameter is a comma separated list of values, so that multiple values can be provided in a single call.

## remote

This `PACKAGE` is the interface for managing the remote connectivity of the current session. This is also the interface to execute remote commands, such as queries, DDL and DML operations.

This object is accessible through the `HIVE_USER` role and is referenced as the `SYNONYM` name `HIVE_REMOTE`

## Interface

### session\_param()

Overloaded `FUNCTION` and `PROCEDURE` to get and set parameters only for the current session. Parameters set using this `PROCEDURE` are valid only for the duration of the Oracle session and are discarded when the session ends.

A session level parameter takes precedence over the system level parameter of the same name.

- Prototype

```
-- get  
function session_param( name in varchar2 ) return varchar2;  
  
-- set  
procedure session_param( name in varchar2,  
                        value in varchar2 );
```

- Parameter

|       |   |  |
|-------|---|--|
| name  | - | Case sensitive name of the parameter to get or <b>set</b>  |
| value | - | <b>Value</b> of the parameter <b>when</b> setting. The <b>value</b> is the <b>return</b> for the <b>FUNCTION</b> |

If the session level parameter is not set, or does not exist then the function returns `NULL`

#### **session\_log\_level()**

A `PROCEDURE` which allows the logging level to be set for the current session. This setting is valid only for the duration of the Oracle session and are discarded when the session ends.

- Prototype

```
procedure session_log_level( typ in number );
```

- Parameter

|            |   |                    |
|------------|---|--------------------|
| <b>typ</b> | - | The log level type |
|------------|---|--------------------|

Log level values are bit masks, to allow for multiple values to be specified in a single numeric value. They are defined as part of the `IMPL` package specification.

```
none constant number := 0;
error constant number := 1;
warn  constant number := 2;
info  constant number := 4;
trace constant number := 8;
```

#### **session\_clear()**

This `PROCEDURE` clears the session `CONNECTION` type.

- Prototype

```
procedure session_clear;
```

After being cleared, consecutive attempts will rebuild the `CONNECTION` type from the parameter information.

#### **session()**

An overloaded `FUNCTION` and `PROCEDURE` which creates, gets, or modifies the current session `CONNECTION` type.

- Prototype

```
-- get
function session return connection;

-- set
procedure session( url in varchar2 );

procedure session( usr in varchar2,
                  pwd in varchar2 );
```

```

procedure session( url in varchar2,
                  usr in varchar2,
                  pwd in varchar2 );

procedure session( url in varchar2,
                  usr in varchar2,
                  pwd in varchar2,
                  ath in varchar2 );

procedure session( con in connection );

```

- Parameter

|     |   |  |
|-----|---|--|
| url | - | The JDBC URL of the remote connection      |
| usr | - | User name to be applied for the connection |
| pwd | - | The user password to be applied            |
| ath | - | The authentication type of the connection  |

When getting the current `CONNECTION` type, the password value is redacted by the function, no matter what its value or the authentication type specified.

#### query()

A `PIPELINED FUNCTION` which returns an `ANYDATASET` value. This value contains the description information of a given row type along with the set of data instances of that row.

- Prototype

```

function query( stm in varchar2,
               bnd in binds default null,
               con in connection default null )
return anydataset pipelined using hive_t;

```

- Parameter

|     |   |  |
|-----|---|--|
| stm | - | The SQL query statement to be executed remotely  |
| bnd | - | Optionally, the <code>HIVE_BINDS</code> array containing the values to be bound to the SQL statement |
| con | - | Optionally, the remote <code>CONNECTION</code> information   |

This function can be used in a casting statement to retrieve the remote records, but it does not contain a `PACKAGE BODY` definition, rather it uses the `HIVE_T` ODCI object type making the following 2 examples equivalent

```

val := hive_q( 'select cust_id,
                last_name,
                first_name
                from cust', null, null );

```

vs.

```
val := hive_remote.query( 'select cust_id,
                           last_name,
                           first_name
                           from cust', null, null );
```

### dml()

This `PROCEDURE` allows remote DML to be executed using the provided SQL statement, bind variables and connection type.

- Prototype

```
procedure dml( stm in varchar2,
               bnd in binds      default null,
               con in connection default null );
```

- Parameter

|            |   |
|------------|---|
| <b>stm</b> | - The DML statement to be executed remotely   |
| <b>bnd</b> | - Optionally, the HIVE_BINDS array containing the values to be bound to the DML statement |
| <b>con</b> | - Optionally, the remote CONNECTION information   |

### ddl()

This `PROCEDURE` allows remote DDL to be executed using the provided command statement and connection type.

- Prototype

```
procedure ddl( stm in varchar2,
               con in connection default null );
```

- Parameter

|            |   |
|------------|---|
| <b>stm</b> | - The DML statement to be executed remotely     |
| <b>con</b> | - Optionally, the remote CONNECTION information |

## impl

This `PACKAGE` is the primary implementation of the ODCI layer and the interface called by Oracle ODCI through the `HIVE_T` object type.

This interface is used internally only, and while it is available for execution with the `HIVE_ADMIN` role it should never be called directly. It is included here only for completeness of the documentation.

There is no SYNONYM for this object and both its specification and body are wrapped during the installation process.

## Constants

## Helper constants for abstracting type information

### Log level

```
none constant number := 0;  
error constant number := 1;  
warn constant number := 2;  
info constant number := 4;  
trace constant number := 8;
```

Log level values are bit masks, to allow for multiple values to be specified in a single numeric value.

## Interface

### log()

This procedure writes a logging record of the type specified, if `log_level` at the systems or session set set to support it.

- Prototype

```
procedure log( typ in number, txt in varchar2 );
```

Additionally, there are pass-through procedures which call this procedure using the predefined type.

```
procedure log_error( txt in varchar2 );  
procedure log_warn( txt in varchar2 );  
procedure log_info( txt in varchar2 );  
procedure log_trace( txt in varchar2 );
```

- Parameter

|     |   |   |
|-----|---|---|
| typ | - | The type <b>of</b> log to write, ignored <b>if</b> the log_level <b>is not</b> set to support that type |
| txt | - | The text <b>of</b> the message  |

### session\_param()

Overloaded `FUNCTION` and `PROCEDURE` to get and set parameters only for the current session. Parameters set using this `PROCEDURE` are valid only for the duration of the Oracle session and are discarded when the session ends.

A session level parameter takes precedence over the system level parameter of the same name.

- Prototype

```
-- get  
function session_param( name in varchar2 ) return varchar2;
```

```
-- set
procedure session_param( name in varchar2,
                        value in varchar2 );
```

- Parameter

|       |  |
|-------|--|
| name  | - Case sensitive name of the parameter to get or <b>set</b>  |
| value | - <b>Value</b> of the parameter <b>when</b> setting. The <b>value</b> is the <b>return</b> for the <b>FUNCTION</b> |

If the session level parameter is not set, or does not exist then the function returns `NULL`

### session\_log\_level()

A `PROCEDURE` which allows the logging level to be set for the current session. This setting is valid only for the duration of the Oracle session and are discarded when the session ends.

- Prototype

```
procedure session_log_level( typ in number );
```

- Parameter

|     |                      |
|-----|----------------------|
| typ | - The log level type |
|-----|----------------------|

### session\_clear()

This `PROCEDURE` clears the session `CONNECTION` type.

- Prototype

```
procedure session_clear;
```

After being cleared, consecutive attempts will rebuild the `CONNECTION` type from the parameter information.

### session()

An overloaded `FUNCTION` and `PROCEDURE` which creates, gets, or modifies the current session `CONNECTION` type.

- Prototype

```
-- get
function session return connection;

-- set
procedure session( url in varchar2 );

procedure session( usr in varchar2,
                  pwd in varchar2 );
```

```

procedure session( url in varchar2,
                  usr in varchar2,
                  pwd in varchar2 );

procedure session( url in varchar2,
                  usr in varchar2,
                  pwd in varchar2,
                  ath in varchar2 );

procedure session( con in connection );

```

- Parameter

|     |   |   |
|-----|---|---|
| url | - | The JDBC URL <b>of</b> the remote connection      |
| usr | - | User name to be applied <b>for</b> the connection |
| pwd | - | The user password to be applied                   |
| ath | - | The authentication type <b>of</b> the connection  |

When getting the current `CONNECTION` type, the password value is **not** redacted by this function, unlike the function of the same name used in the `REMOTE` package

#### sql\_describe()

This overloaded `FUNCTION` is called by ODCI to describe information for a table whose return type is `ANYDATASET`. This function is called when `HIVE_T` `ODCITableDescribe` is encountered.

- Prototype

```

function sql_describe( stm in varchar2,
                      bnd in binds      default null,
                      con in connection default null ) return anytype;

function sql_describe( typ out anytype,
                      stm in  varchar2,
                      bnd in  binds      default null,
                      con in  connection default null ) return number;

function sql_describe( key in  number,
                      typ out anytype ) return number;

```

- Parameter

|     |   |   |
|-----|---|---|
| stm | - | The SQL query statement to be executed remotely   |
| bnd | - | Optionally, the HIVE_BINDS array containing the values to be bound to the SQL statement                           |
| con | - | Optionally, the remote CONNECTION information   |
| typ | - | The ANYTYPE <b>value</b> that describes the returned rows <b>from</b> the table function                          |
| key | - | The transient key <b>value</b> used <b>in</b> the concurrent calling chain, identifying the record of the context |

If successful the return value for each function will be

`ODCICONST.SUCCESS` otherwise it will return `ODCICONST.ERROR`

### **sql\_open()**

This `FUNCTION` is called by ODCI to initialize the scan of a table function. This function is called when `HIVE_T` `ODCITableStart` is encountered.

- Prototype

```
function sql_open( key out number,  
                  stm in varchar2,  
                  bnd in binds      default null,  
                  con in connection default null ) return number;
```

- Parameter

|     |   |   |
|-----|---|---|
| key | - | The transient key value used <b>in</b> the concurrent calling chain, identifying the record <b>of</b> the context |
| stm | - | The SQL query statement to be executed remotely   |
| bnd | - | Optionally, the HIVE_BINDS array containing the values to be bound to the SQL statement                           |
| con | - | Optionally, the remote CONNECTION information   |

If successful the return value for each function will be

`ODCICONST.SUCCESS` otherwise it will return `ODCICONST.ERROR`

### **sql\_fetch()**

This `FUNCTION` is called by ODCI to next batch of rows from a table function. This function is called when `HIVE_T` `ODCITableFetch` is encountered.

- Prototype

```
function sql_fetch( key in number,  
                   num in number,  
                   rws out records ) return number;
```

- Parameter

|     |   |   |
|-----|---|---|
| key | - | The transient key value used <b>in</b> the concurrent calling chain, identifying the record <b>of</b> the context |
| num | - | The number <b>of</b> rows the system expects <b>in</b> the current fetch cycle.                                   |
| rws | - | The RECORDS array <b>for</b> the rows requested <b>in</b> the current fetch cycle.                                |

If successful the return value for each function will be

`ODCICONST.SUCCESS` otherwise it will return `ODCICONST.ERROR`

### **sql\_close()**

This `FUNCTION` is called by ODCI to performs cleanup operations



after scanning a table function cycle is complete. This function is called when `HIVE_T` `ODCITableClose` is encountered.

- **Prototype**

```
function sql_close( key in number ) return number;
```

- **Parameter**

|     |   |   |
|-----|---|---|
| key | - | The transient key value used in the concurrent calling chain, identifying the record of the context |
|-----|---|---|

If successful the return value for each function will be `ODCICONST.SUCCESS` otherwise it will return `ODCICONST.ERROR`

### **sql\_dml()**

This `PROCEDURE` allows remote DML to be executed using the provided SQL statement, bind variables and connection type.

- **Prototype**

```
procedure sql_dml( stm in varchar2,  
                  bnd in binds      default null,  
                  con in connection default null );
```

- **Parameter**

|     |   |   |
|-----|---|---|
| stm | - | The DML statement to be executed remotely   |
| bnd | - | Optionally, the HIVE_BINDS array containing the values to be bound to the DML statement |
| con | - | Optionally, the remote CONNECTION information   |

### **sql\_ddl()**

This `PROCEDURE` allows remote DDL to be executed using the provided command statement and connection type.

- **Prototype**

```
procedure sql_ddl( stm in varchar2,  
                  con in connection default null );
```

- **Parameter**

|     |   |   |
|-----|---|---|
| stm | - | The DML statement to be executed remotely     |
| con | - | Optionally, the remote CONNECTION information |

## **Exception**

### **ex\_not\_eligible**

This exception is thrown when a request to change a parameter at the session level is ineligible for the operation.

```
ex_not_eligible exception;
es_not_eligible constant varchar2( 256 ) := 'Parameter is not
                                             eligible for
                                             change at the
                                             session level';

ec_not_eligible constant number           := -20103;
```

## Functions

Only the `HIVE_Q` function is available for execution, while all other functions are for internal use only.

### hive\_q

A `PIPELINED FUNCTION` which returns an `ANYDATASET` value. This value contains the description information of a given row type along with the set of data instances of that row.

This function is accessible through the `HIVE_USER` role and is referenced as the `SYNONYM` name `HIVE_Q`

- Prototype

```
function hive_q( stm in varchar2,
                 bnd in binds      default null,
                 con in connection default null )
return anydataset pipelined using hive_t;
```

- Parameter

|            |   |
|------------|---|
| <b>stm</b> | - The SQL query statement to be executed remotely   |
| <b>bnd</b> | - Optionally, the HIVE_BINDS array containing the values to be bound to the SQL statement |
| <b>con</b> | - Optionally, the remote CONNECTION information   |

This function can be used in a casting statement to retrieve the remote records, but it does not contain a `PACKAGE BODY` definition, rather it uses the `HIVE_T` ODCI object type making the following 2 examples equivalent

```
val := hive_q( 'select cust_id,
               last_name,
               first_name
               from cust', null, null );
```

vs.

```
val := hive_remote.query( 'select cust_id,
```

```
last_name,
first_name
from cust', null, null );
```

## hive\_hint

This helper `FUNCTION`, accessible through the `HIVE_USER` role, returns a hints for an `OWNER.TABLE` column listing. The hint string returned can be passed into a select for a `HIVE_Q` query. Please see *Hints* section above for more information

- Prototype

```
function hive_hint( own in varchar2,
                   tab in varchar2,
                   typ in varchar2 default 'typecast' )
```

- Parameter

|     |   |   |
|-----|---|---|
| own | - | Owner of the <b>template</b> table to generate hints from |
| tab | - | Table <b>template</b> to generate hints from              |
| typ | - | Type of hint <b>string</b> to generate                    |

## bitxor

This `FUNCTION` performs the logical exclusive OR operation on each pair of corresponding bits provided.

- Prototype

```
function bitxor( x in number, y in number ) return number;
```

- Parameter

|   |   |                 |
|---|---|-----------------|
| x | - | First bit mask  |
| y | - | Second bit mask |

## bitnot

This `FUNCTION` performs logical negation on each bit, forming the complement of the given binary value provided.

- Prototype

```
function bitnot( x in number ) return number;
```

- Parameter

|   |   |          |
|---|---|----------|
| x | - | Bit mask |
|---|---|----------|

## oid

This `FUNCTION` returns the database object identifier for the name provided, which is expected to be a `USERNAME` or `ROLE`

The return value is used as the join key for the saved Filter privileges

- Prototype

```
function oid( o in varchar2 ) return number;
```

- Parameter

`o` - Object name

## oname

This `FUNCTION` returns the database object name for the identifier provided, which is expected to be a `USER#` representing the `USERNAME` or `ROLE`

The return value is used as the join key displaying names for the saved Filter privileges.

- Prototype

```
function oname( o in number ) return varchar2;
```

- Parameter

`o` - Object identifier

## Types

---

Hive-ODCI types are both objects and table (arrays) and are used in the transparent conversions from the JDBC Driver to the PL/SQL calls.

## attribute

This `OBJECT` type is used to describe a column value. it is used internally only to define rows in a record set and a table description.

- Members

|                    |                                       |
|--------------------|---------------------------------------|
| <code>name</code>  | - Name <b>of</b> the column           |
| <code>code</code>  | - Data type code <b>of</b> the column |
| <code>prec</code>  | - The precision <b>of</b> the column  |
| <code>scale</code> | - The scale <b>of</b> the column      |
| <code>len</code>   | - The length <b>of</b> the column     |
| <code>csid</code>  | - The locale identifier               |

csfrm - The locale format

## attributes

A **TABLE** array of **ATTRIBUTE** objects.

- Members

table **of** attribute

## data

This **OBJECT** type contains a single column of information.

- Members

|               |   |
|---------------|---|
| code          | - The data type code, pointing to the member which actually contains the data |
| val_varchar2  | - <b>String</b> data  |
| val_number    | - Numeric data  |
| val_date      | - <b>Date</b> data  |
| val_timestamp | - Timestamp data  |
| val_clob      | - CLOB data   |
| val_blob      | - BLOB data   |

## records

A **TABLE** array of **DATA** objects, representing a single row

- Members

table **of** data

## connection

This **OBJECT** type contains the remote connection information

- Members

|      |   |
|------|---|
| url  | - The JDBC URL <b>of</b> the remote connection      |
| name | - User name to be applied <b>for</b> the connection |
| pass | - The user password to be applied                   |
| auth | - The authentication type <b>of</b> the connection  |

## bind

This **OBJECT** defines the bind variable used for the remote command

- Members

|             |   |
|-------------|---|
| value       | - The value of the data to be bound                                 |
| <b>type</b> | - The data <b>type</b> code, <b>for</b> how to <b>bind</b> the data |

scope - The reference code direction of the [bind](#)

## binds

A `TABLE` array of `BIND` objects

- Members

table of bind

## hive\_t

An `OBJECT` type for the data cartridge abstraction later. This object is the compliant interface for ODCI callback functionality.

- Members

|     |   |
|-----|---|
| key | - The transient key <b>value</b> used <b>in</b> the concurrent calling chain, identifying the record of the context |
| ref | - The persistent ANYTYPE reference being populated <b>in</b> the current calling cycle                              |

## ODCITableDescribe

Retrieves describe information for a table function whose return type is `ANYDATASET`.

- Prototype

```
function ODCITableDescribe( typ out anytype,  
                           stm in varchar2,  
                           bnd in binds      := null,  
                           con in connection := null ) return number
```

- Parameter

|     |  |
|-----|--|
| typ | - The AnyType <b>value</b> that describes the returned rows <b>from</b> the table function |
| stm | - The command statement to be executed remotely  |
| bnd | - Optionally, the HIVE_BINDS array containing the values to be bound to the statement      |
| con | - Optionally, the remote CONNECTION information  |

## ODCITablePrepare

This `FUNCTION` prepares the scan context and command information upon a compile request.

- Prototype

```
function ODCITablePrepare( ctx out hive_t,  
                          inf in sys.ODCITabFuncInfo,  
                          stm in varchar2,
```

```
bnd in binds      := null,
con in connection := null ) return number
```

- Parameter

|     |   |  |
|-----|---|--|
| ctx | - | The scan context created by <b>this</b> routine <b>is</b> the <b>value</b> passed <b>in as</b> a parameter to the later routines <b>in</b> the command cycle |
| inf | - | The projection information and the <b>return</b> descriptor <b>object</b>  |
| stm | - | The command statement to be executed remotely  |
| bnd | - | Optionally, the HIVE_BINDS array containing the values to be bound to the statement  |
| con | - | Optionally, the remote CONNECTION information  |

## ODCITableStart

This **FUNCTION** initializes the scan of a table function to start the command cycle.

- Prototype

```
function ODCITableStart( ctx in out hive_t,
                        stm in      varchar2,
                        bnd in      binds      := null,
                        con in      connection := null ) return number
```

- Parameter

|     |   |   |
|-----|---|---|
| ctx | - | The scan context modified by <b>this</b> routine <b>is</b> the <b>value</b> passed <b>in as</b> a parameter to the later routines <b>in</b> the command cycle |
| stm | - | The command statement to be executed remotely   |
| bnd | - | Optionally, the HIVE_BINDS array containing the values to be bound to the statement   |
| con | - | Optionally, the remote CONNECTION information   |

## ODCITableFetch

This **FUNCTION** returns the next batch of rows in the command cycle.

- Prototype

```
function ODCITableFetch( cts in out hive_t,
                        num in      number,
                        rws out     anydataset ) return number
```

- Parameter

|     |   |  |
|-----|---|--|
| ctx | - | The scan context modified <b>by this</b> routine <b>is</b> the value passed <b>in</b> as a parameter to the later routines <b>in</b> the command cycle |
| num | - | The number <b>of</b> rows the system expects <b>in</b> the current   |

```
    rws      -   fetch cycle.  
                  The RECORDS array for the rows requested in the  
                  current fetch cycle.
```

## ODCITableClose

This `FUNCTION` performs cleanup operations after command cycle is complete

- Prototype

```
function ODCITableClose( ctx in hive_t ) return number
```

- Parameter

```
ctx      -   The scan context created by this routine is the value  
              passed in as a parameter to the later routines  
              in the command cycle
```

## Parameters

---

Parameters provide the customization of Hive-ODCI.

### **application**

The application name

### **version**

The installed or patched version

### **license**

The BSD licensing agreement

### **log\_level**

The current log level

### **hive\_jdbc\_driver**

The fully qualified Java class name of the driver to load

### **hive\_jdbc\_url**

The JDBC URL for the remote system

### **hive\_jdbc\_url.X**

These parameters are additional URL key value pairs in consecutive order. A gap in the numbering sequence will cause Hive-ODCI to stop reading the parameters assuming that it has reached the end of the list.



## java\_property.X

These parameters are the Java Systems properties to set in consecutive order. A gap in the numbering sequence will cause Hive-ODCI to stop reading the parameters assuming that it has reached the end of the list.

## hive\_user

The `Driver.getConnection()` call in Java is overloaded to accept a URL only or optionally with a User/Password. When set this parameter is passed through to the call as the user value.

## hive\_pass

The `Driver.getConnection()` call in Java is overloaded to accept a URL only or optionally with a User/Password. When set this parameter is passed through to the call as the password value.

## hive\_auth

The authentication type for the JDBC Driver.

## query\_limit

The ceiling limitation for any given query,. When set no query will return more rows than the parameter specified

## auth\_no\_grant\_action

Determine rule when no table grants are available

options:

- `error` Throw an exception to the calling process (log is implied)
- `log` Log the this is an error, but take no further action
- `auto` Add new grants, based on the "auth\_grant\_handler" parameter rules
- `ignore` Do nothing, take no action

default:

`ignore`

## auth\_auto\_grant

Define the grant handling rules, when auto is set in `auth_no_grant_action` this is a comma separate list of users and/or roles

options:

- `current` A special-case name, meaning the currently logged in user
- `<role>` Any valid role name (invalids ignored)
- `<user>` Any valid user name (invalids ignored)

default:

`NULL` (a `NULL` or missing value, takes no action)

## auth\_table\_undefined

Determine rule when no table grants are found

options:

- `error` Throw an exception to the calling process ( `log` is implied)
- `allow` Assume allow access
- `deny` Assume deny access

default:

`allow`

## auth\_sql\_parse\_error

Determine rule when handling JSqlParser error (parsing sql or syntax error)

option:

- `error` Throw the JSqlParserException back to the calling process
- `none` Catch the error and return none (e.g. no table info)

default:

`none`

## Roles

---

### hive\_user

This role provides the access privileges necessary for common Hive-ODCI functionality.

### hive\_admin

This role provides the escalated access privileges necessary used by administrators of Hive-ODCI.

## Synonyms

---

Hive-ODCI synonyms are the public alternative names for the interface objects providing location transparency.

- `hive_q`
- `hive_t`
- `hive_remote`
- `hive_bind`
- `hive_binds`
- `hive_binding`
- `hive_attribute`
- `hive_attributes`
- `hive_data`
- `hive_records`
- `hive_connection`
- `dbms_hive`
- `dba_hive_params`

- dba\_hive\_filters
- dba\_hive\_filter\_privs
- dba\_hive\_privs
- dba\_hive\_log
- user\_hive\_params
- user\_hive\_filters
- user\_hive\_filter\_privs
- user\_hive\_privs